



# Tools for Artificial Intelligence with MATLAB, initiation (TAIM)

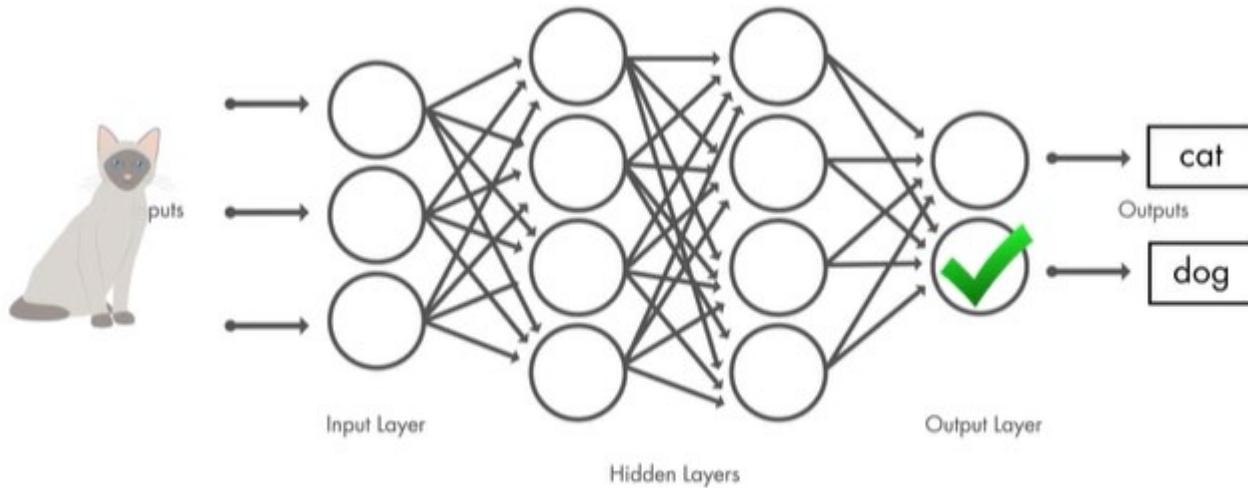
**José Antonio Lázaro**

**Deep Learning**

Barcelona, 4, February 2025

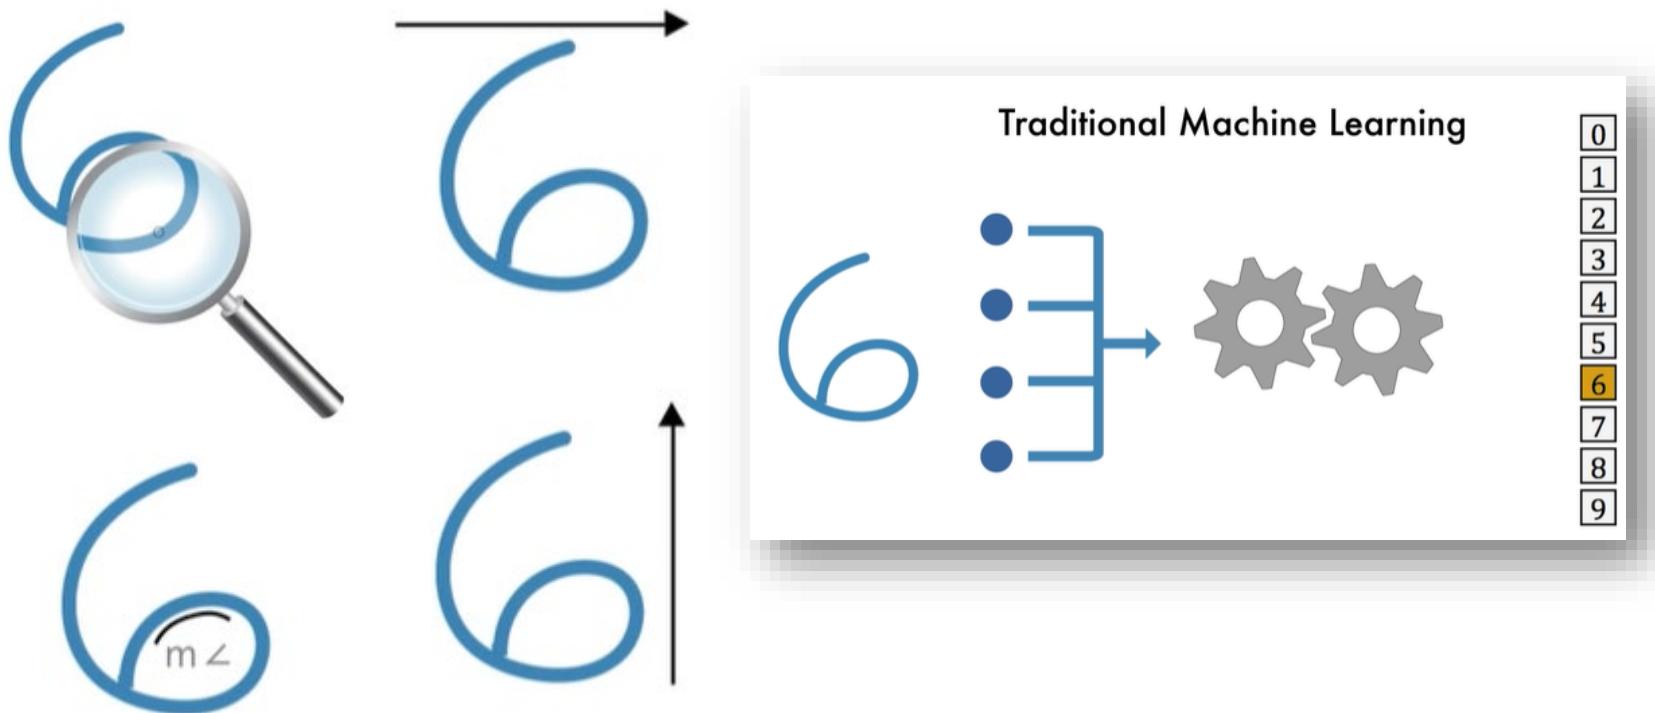
# What is it?

- Deep learning is a Machine learning technique that uses Deep Neural Networks (networks with many layers) to make predictions.



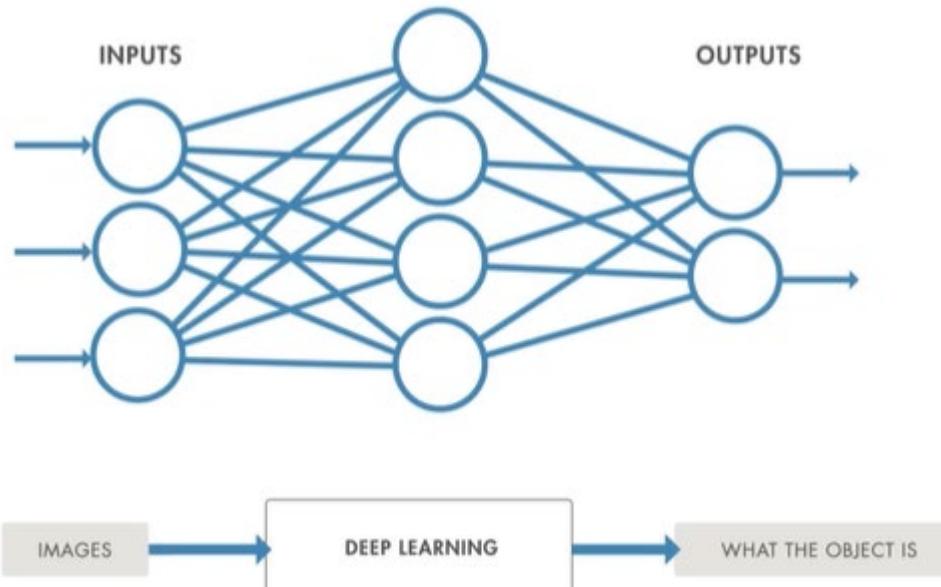
# Comparisson with Machine Learning

- Deep learning is a Machine learning technique that uses Deep Neural Networks (networks with many layers) to make predictions.



# Comparisson with Machine Learning

- Deep learning goal is to perform end-to-end learning.
- The images themselves are the inputs.
- Features & Classification are learned directly from the images.

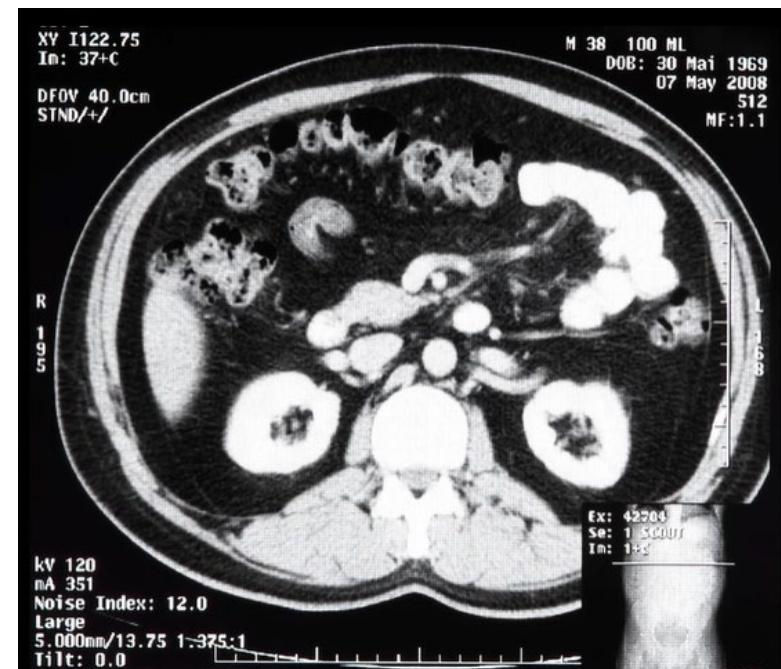
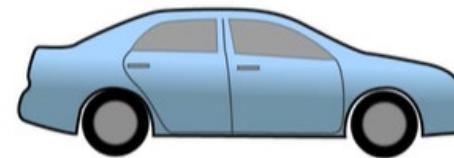


# What for?

- Extracting information from Image.

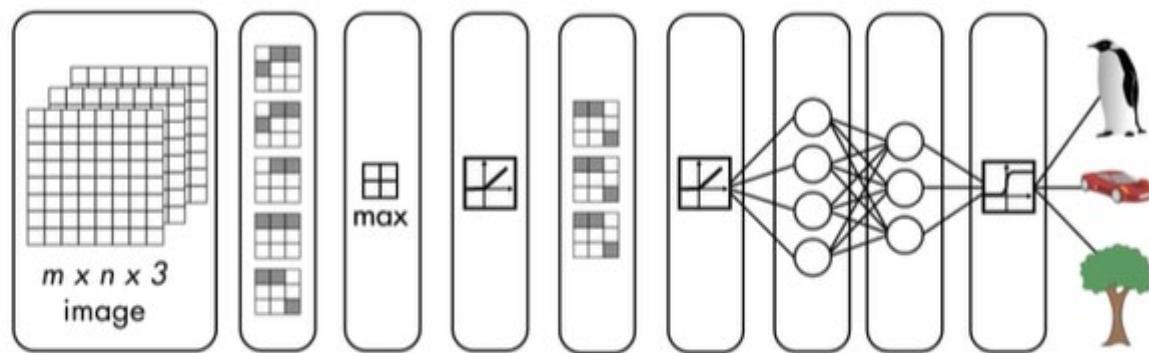
E.g.:

- Hand-written text recognition
- Autonomous Vehicles
- Medical preliminary diagnosis from the image



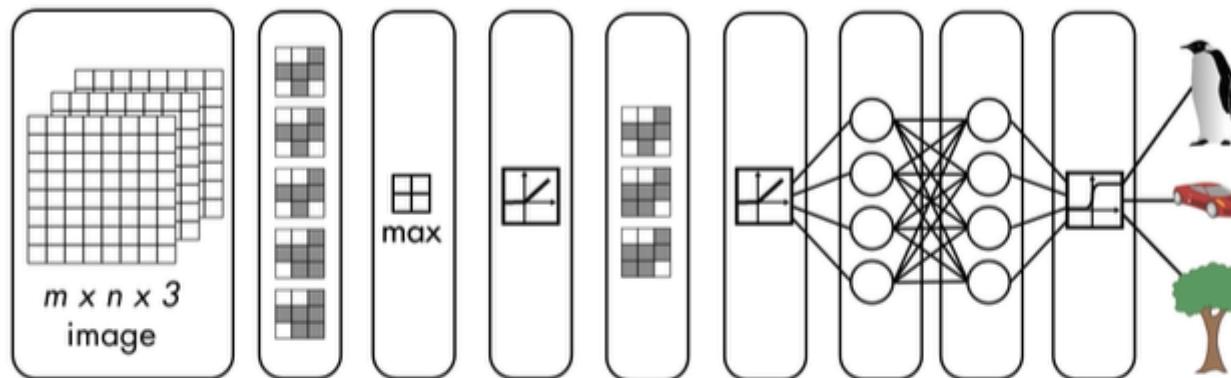
# General Scheme

- Extracting information from data.



# General Scheme

- As in previous applications, Matlab will help implementing the mathematical formulas.



$$E(\theta) = - \sum_{i=1}^n \sum_{j=1}^k t_{ij} \ln y_j(x_i, \theta)$$

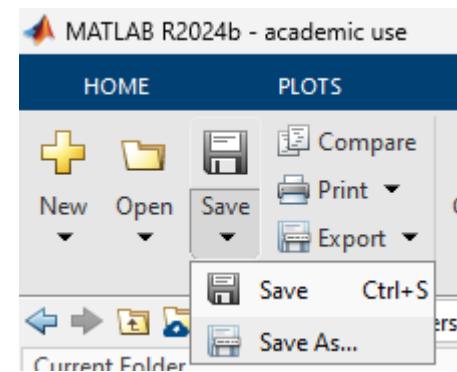
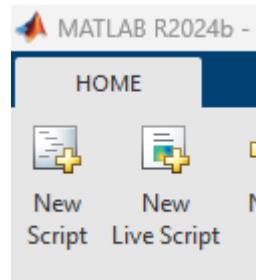
$$\theta_{\ell+1} = \theta_\ell - \alpha \nabla E(\theta_\ell) + \gamma (\theta_\ell - \theta_{\ell-1})$$

$$y_r(x) = \frac{\exp(a_r(x))}{\sum_{j=1}^k \exp(a_j(x))}$$

# Let's go

## Open Matlab

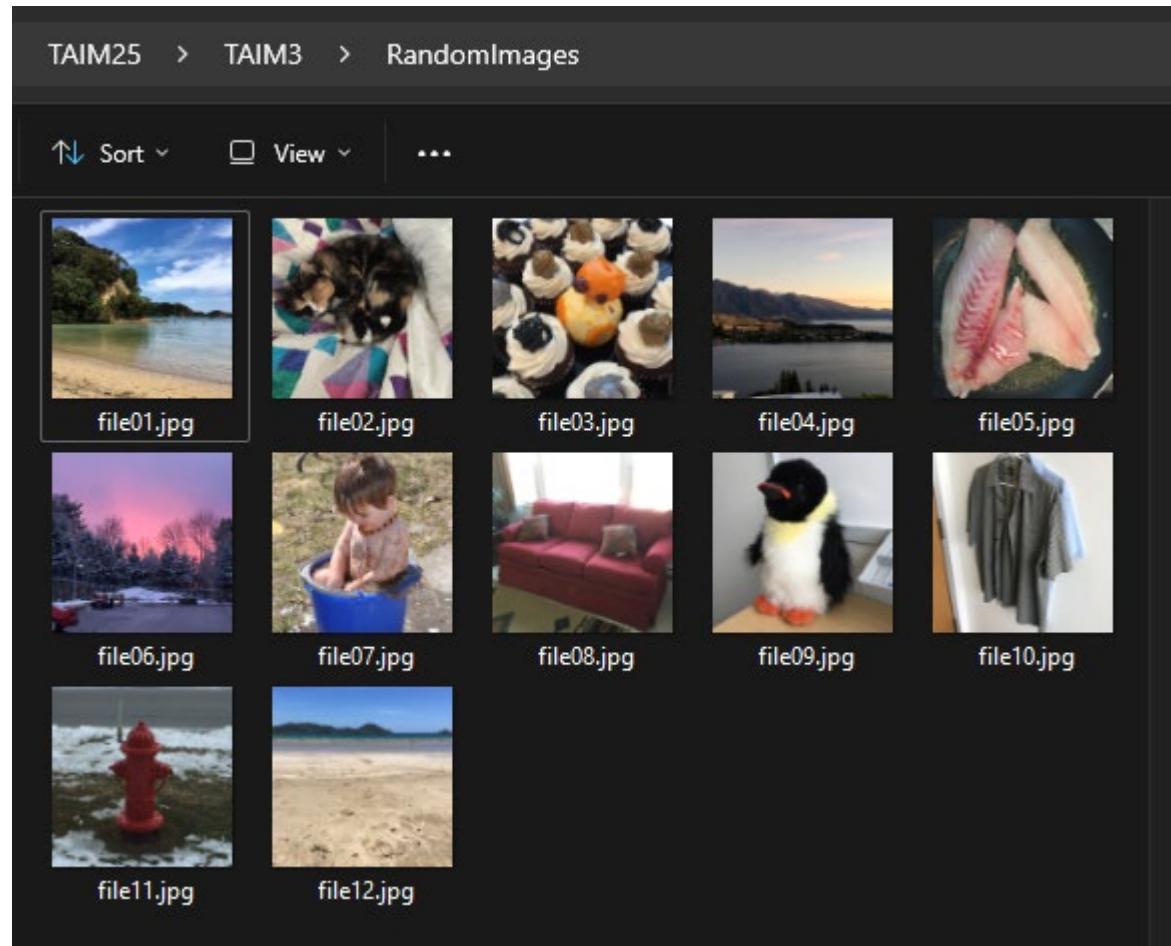
- Do “New Live Script”
- Go “Live Editor”
- And “Save as”
- Create a new Folder (E.g. “**TAIM3**”)
- NO SPACES at the NAME



TAIM25 > TAIM3 >	
↑ Sort ▾	
Name	Date modified
Others	2/5/2025 6:04 PM
RandomImages	2/5/2025 6:03 PM
Roundworms	2/5/2025 6:03 PM
LICENSE.txt	2/5/2025 6:02 PM
TAIM_3.mlx	2/5/2025 7:01 PM

# Let's start Preparing DATA

- Download and paste to your “TAIM3” the unzipped folder  
**“RandomImages”**



# Let's start Preparing DATA

- Download and paste to your “TAIM3” the unzipped folder  
**“RandomImages”**

All images in this archive are licensed under the Creative Commons By-Attribution License, available at:  
<https://creativecommons.org/licenses/by/2.0/>  
The photographers are listed below, thanks to all of them for making their work available, and please be sure to credit them for any use as per the license.

```
daisy/7568630428_8cf0fc16ff_n.jpg CC-BY by A Guy Taking Pictures - https://www.flickr.com/photos/80901381@N04/7568630428/
daisy/7410356270_9dff4d0e2e_n.jpg CC-BY by martinak15 - https://www.flickr.com/photos/martinaphotography/7410356270/
daisy/4286053334_a75541f20b_m.jpg CC-BY by jenny downing - https://www.flickr.com/photos/jenny-pics/4286053334/
daisy/10770585085_4742b9dac3_n.jpg CC-BY by Thangaraj Kumaravel - https://www.flickr.com/photos/kumaravel1/10770585085/
daisy/8759177308_951790e00d_m.jpg CC-BY by Marina del Castell - https://www.flickr.com/photos/marinadelcastell/8759177308/
daisy/4131565290_0585c4dd5a_n.jpg CC-BY by jenny downing - https://www.flickr.com/photos/jenny-pics/4131565290/
daisy/3504430338_77d6a7fab4_n.jpg CC-BY by Dhilung Kirat - https://www.flickr.com/photos/dhilung/3504430338/
daisy/3084924076_4d5c5711af_m.jpg CC-BY by Büi Linh Ngân - https://www.flickr.com/photos/linhngan/3084924076/
daisy/2642408410_61545fdc83_n.jpg CC-BY by Dennis Jarvis - https://www.flickr.com/photos/archer10/2642408410/
daisy/8710109684_e2c5ef6aeb_n.jpg CC-BY by Jon Bunting - https://www.flickr.com/photos/84744710@N06/8710109684/
daisy/2612704455_efce1c2144_m.jpg CC-BY by Swaminathan - https://www.flickr.com/photos/araswami/2612704455/
daisy/8021540573_c56cf9070d_n.jpg CC-BY by martinak15 - https://www.flickr.com/photos/martinaphotography/8021540573/
daisy/4412840840_184262b51_n.jpg CC-BY by Neal Evans - https://www.flickr.com/photos/21878512@N06/4412840840/
```

licensed under the Creative Commons By-Attribution License, available at:  
<https://creativecommons.org/licenses/by/2.0/>

# Let's start coding

- Let's load and see images by:

```
"img1 =  
imread("file01.jpg")  
  
imshow(img1)"
```

How you would show some other images? → Do it!

```
img1 = 227×227×3 uint8 array  
img1(:,:,1) =
```

90	90	89	87	85	84	83	82	79	80	82	85	89
91	91	90	88	87	85	84	84	80	81	82	84	88
93	92	91	90	89	88	87	87	82	82	82	83	85
95	94	94	93	92	91	91	90	86	85	84	84	85
97	97	96	96	95	95	95	94	92	91	87	86	88
99	99	99	99	98	98	98	98	96	95	94	91	91
100	100	101	101	101	101	101	101	101	99	98	98	95
101	101	102	102	102	102	102	103	103	103	101	101	98
103	102	115	105	106	113	108	104	94	95	96	97	99
112	100	105	108	109	110	112	96	98	111	105	102	104
120	132	120	120	112	102	96	123	120	94	107	103	109
128	90	118	119	117	109	51	84	85	63	95	63	81
119	100	128	117	124	120	75	64	59	81	79	72	60
134	145	67	77	82	62	71	89	110	92	73	48	50



# Using Neural Networks

- Starting “from scratch” is hard.
- Pretrained networks have been trained by experts to classify images into hundreds of categories.
- We will 1st load a pretrained network into MATLAB and then use the network with new images.

```
% Create a copy of a pretrained deep network
```

```
net = imagePretrainedNetwork
```

```
net =  
dlnetwork with properties:
```

```
Layers: [68x1 nnet.cnn.layer.Layer]  
Connections: [75x2 table]  
Learnables: [52x3 table]  
State: [0x3 table]  
InputNames: {'data'}  
OutputNames: {'prob_flatten'}  
Initialized: 1
```

View summary with `summary`.

```
% By default, the function provides a deep network named  
“dlnetwork”.
```

```
% This network was trained on a data set called ImageNet,  
% which contains 1000 classes, or categories.
```

# Using Neural Networks

**Let's ask this “dlnetwork” to predict file01 to file03 and see the result.**

% You can use the minibatchpredict function to predict the subject of an image.

% The result is a vector of prediction scores, one for each class that the network net can predict.

```
% scores = minibatchpredict(net,img)
```

```
s1 = minibatchpredict(net,img1)
```

*s1 = 1x1000 single row vector*

```
0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 0.0000 ...
```



# Using Neural Networks

% plot(s) to see what do we have obtained

% How do we know what means the result?

% Yes, we forgot to "load" the "classNames" (1000 classes, or categories)

% of "dlnetwork"

% We can do by:

[net,classes] = imagePretrainedNetwork

net =

dlnetwork with properties:

Layers: [68x1 nnet.cnn.layer.Layer]

Connections: [75x2 table]

Learnables: [52x3 table]

State: [0x3 table]

InputNames: {'data'}

OutputNames: {'prob\_flatten'}

Initialized: 1

View summary with summary.

classes = 1000x1 string

"tench"

"goldfish"

"great white s..."

"tiger shark"

"hammerhead"

"electric ray"

"stingray"

"cock"

"hen"

"ostrich"

# Using Neural Networks

% Now the file01 is:

```
c1 = scores2label(s1,classes)
```

c1 = categorical  
seashore

% Now to predict file02 and file03 and see the result:

```
img2 = imread("file02.jpg");  
imshow(img2)
```

```
s2 = minibatchpredict(net,img2)
```

```
c2 = scores2label(s2,classes)
```

c2 = categorical  
Gordon setter



# Using Neural Networks

```
img3 = imread("file03.jpg");  
imshow(img3)
```

```
s3 = minibatchpredict(net,img3)
```

```
c3 = scores2label(s3,classes)
```

c3 = categorical  
bakery



# Other Neural Networks?

Let's try another network to predict file01 to file03 and see the result (E.g. "GoogLeNet")

[net,classes] = imagePretrainedNetwork("googlenet")

net =

dlnetwork with properties:

Layers: [143x1 nnet.cnn.layer.Layer]

Connections: [169x2 table]

Learnables: [116x3 table]

State: [0x3 table]

InputNames: {'data'}

OutputNames: {'prob'}

Initialized: 1

View summary with summary.

classes = 1000x1 string

"tench"

"goldfish"

"great white s..."

"tiger shark"

"hammerhead"

"electric ray"

"stingray"

"cock"

"hen"

"ostrich"

% GoogLeNet has the same prediction

% for the seashore and couch, but predicts that the cat image  
is a different dog breed.

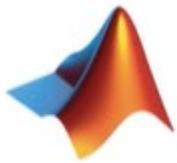
% You can try different pretrained networks

% by looking at the model name arguments in the

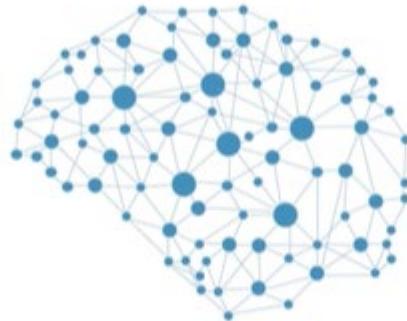
imagePretrainedNetwork

% documentation

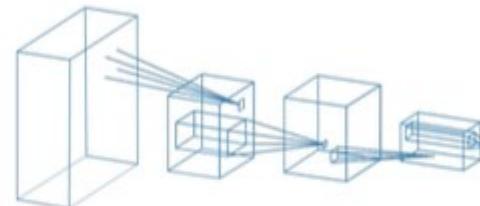
# Other Neural Networks?



MATLAB



Deep Learning  
Toolbox

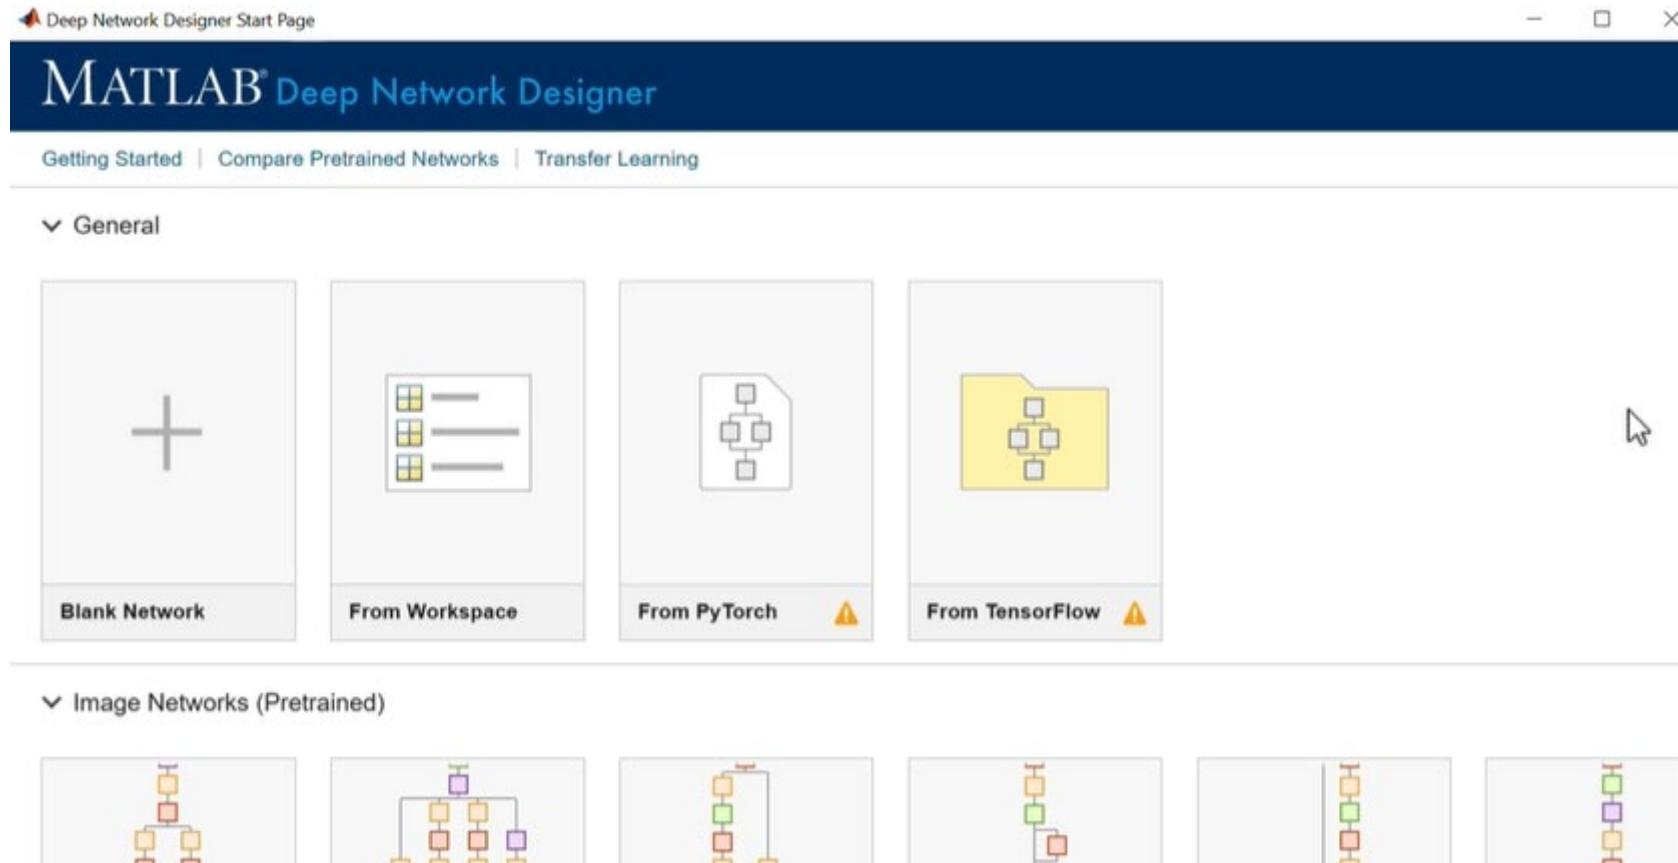


Pretrained  
Network

- Matlab provides access to many of them
- In local Version, you will need to install most of them
- In **Matlab Online**, you have them all preinstalled.

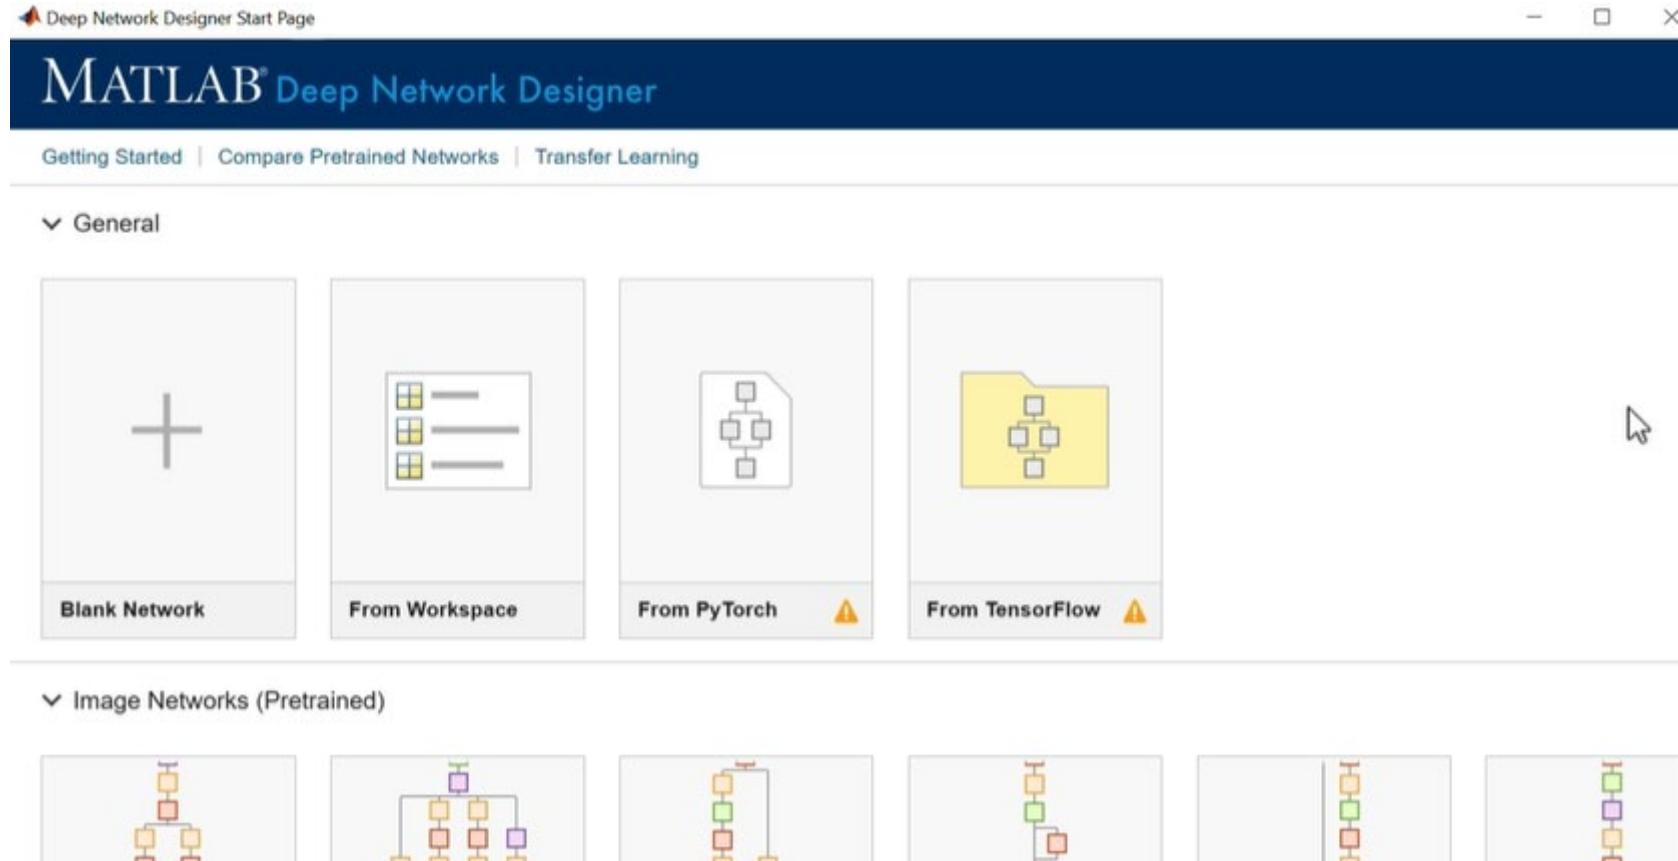
# Yes we have also a Matlab App: Classification Learner

- Following similar approach as with the Curve Fitter App



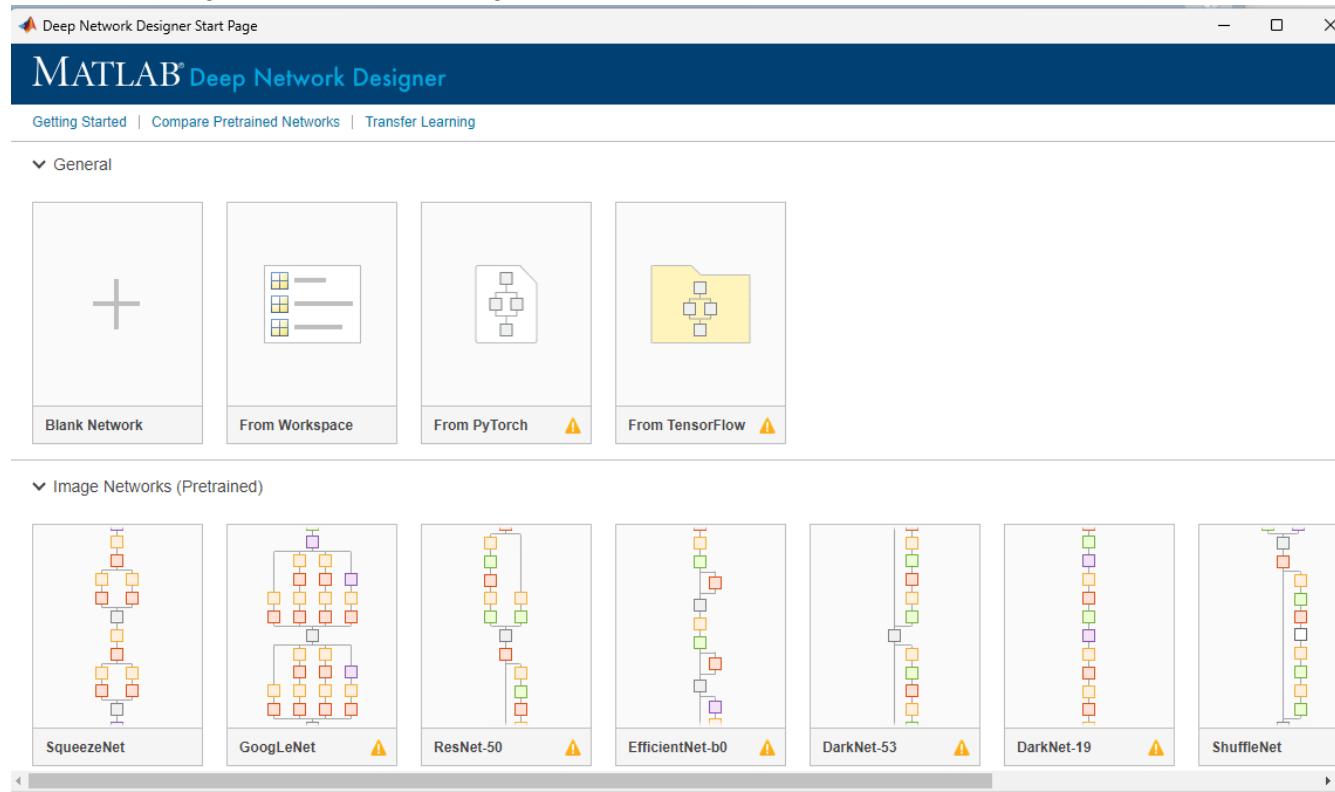
# Yes we have also a Matlab App: Classification Learner

- Following similar approach as with the Curve Fitter App



# Yes we have also a Matlab App: Classification Learner

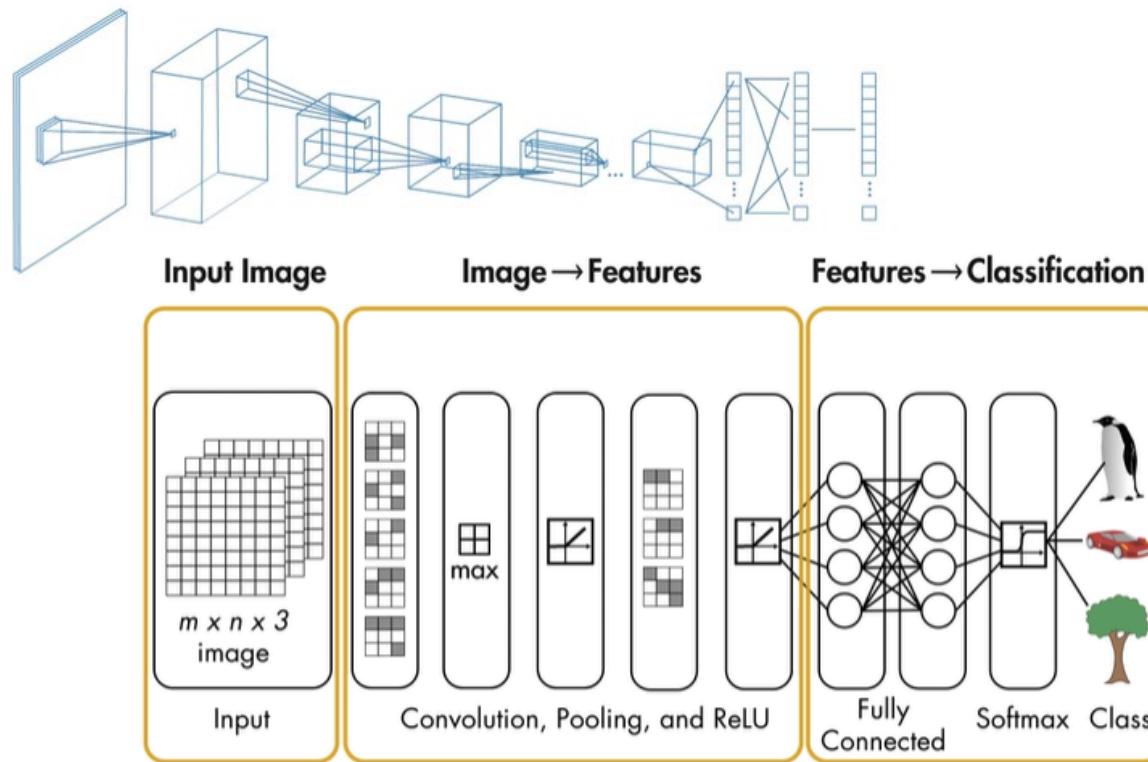
- SqueezeNet is included with your installation, so it can be used right away. Networks with the warning icon need to be installed, so click on the network you want to try. This opens the **Add-On Explorer**.



- To avoid multiple installations, we will use the **Matlab Online**

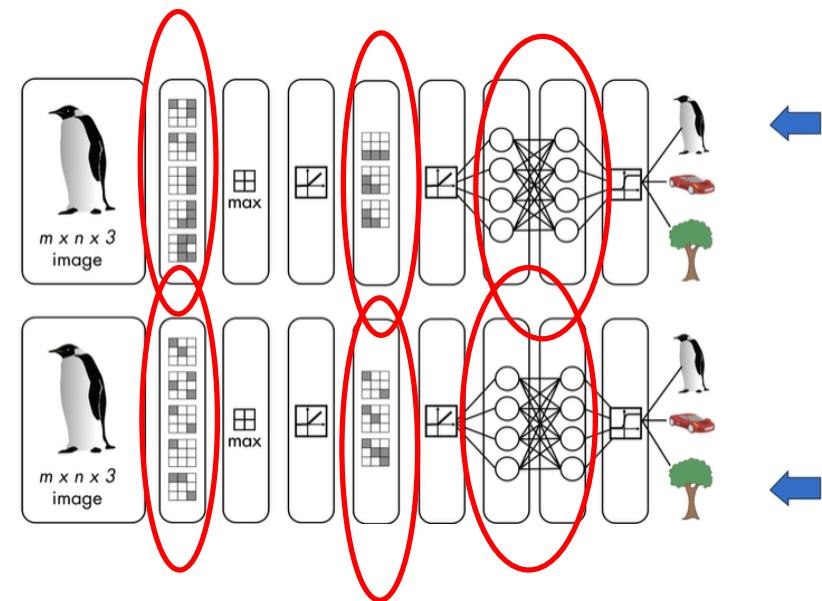
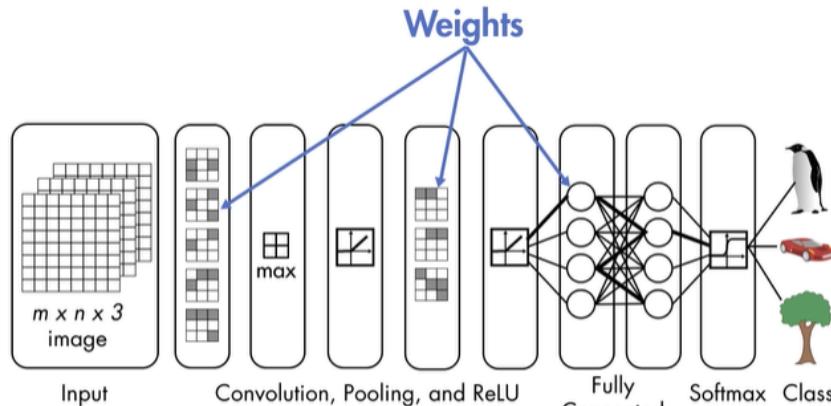
# Convolutional Neural Network

- CNNs are specialized neural networks for inputs with a two-dimensional structure (as images)
- Network architecture:
  - Various layers (key part of the network architect)
  - Layers contain parameters (“weights”)
  - The “weights” are learned by training on labeled data.
- Output layer of a CNN provides predictions for each possible class.



# Convolutional Neural Network

- CNNs are specialized neural networks for inputs with a two-dimensional structure (as images)
- Network architecture:
  - Various layers (key part of the network architect)
  - Layers contain parameters (“weights”)
  - The “weights” are learned by training on labeled data.
- Output layer of a CNN provides predictions for each possible class.



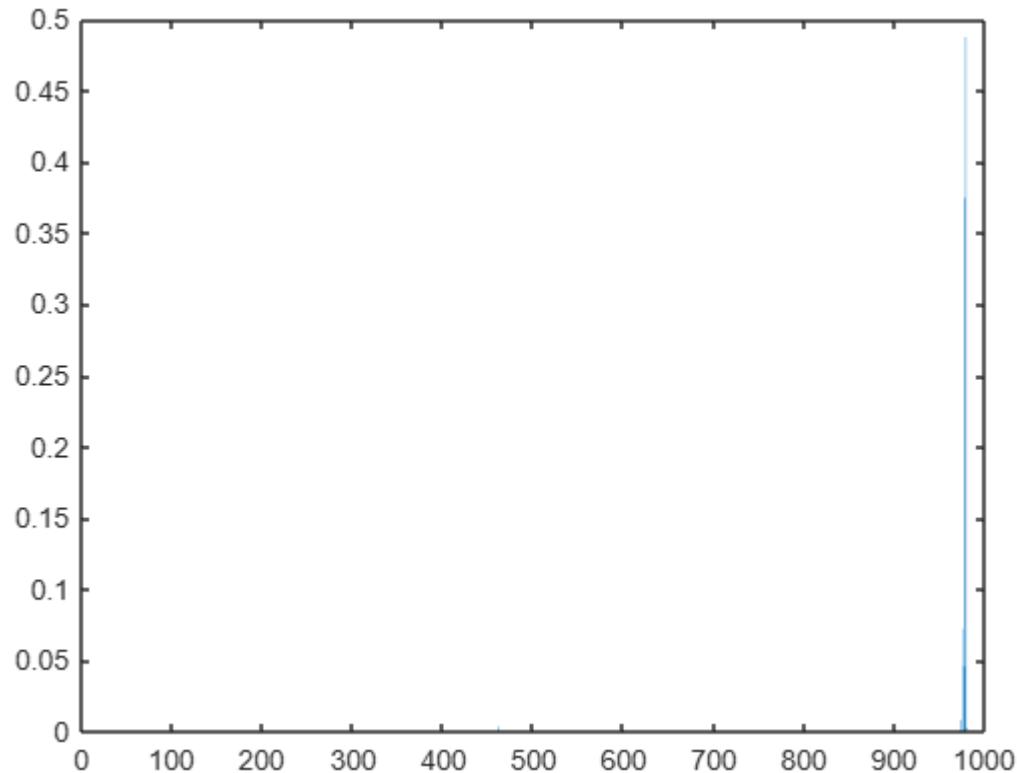
# Examining Predictions

```
% Let's analize predictions for the file01.jpg
```

```
scores = minibatchpredict(net,img1)
```

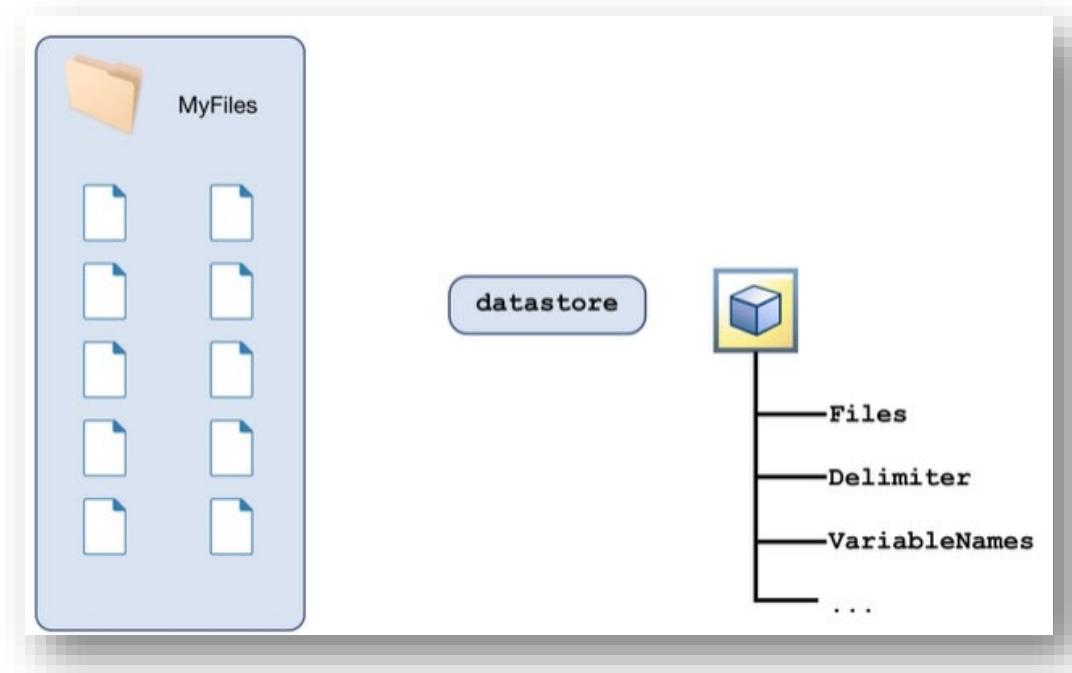
```
% Let's draw a "bar" plot
```

```
bar(scores)
```



Let's go for a “datastore” as in ML

Now is your turn to do it for **file02.jpg** and **file03.jpg**



# Create a Datastore for images

```
ls ("./RandomImages/*.jpg")
```

```
./RandomImages/file01.jpg ./RandomImages/file05.jpg ./RandomImages/file09.jpg ./RandomImages/file02.jpg  
./RandomImages/file06.jpg ./RandomImages/file10.jpg ./RandomImages/file03.jpg ./RandomImages/file07.jpg  
./RandomImages/file11.jpg ./RandomImages/file04.jpg ./RandomImages/file08.jpg ./RandomImages/file12.jpg
```

```
imds = imageDatastore("./RandomImages/*.jpg")
```

```
imds =  
  ImageDatastore with properties:  
  
    Files: {  
      '/MATLAB Drive/TAIM25/TAIM3/RandomImages/file01.jpg';  
      '/MATLAB Drive/TAIM25/TAIM3/RandomImages/file02.jpg';  
      '/MATLAB Drive/TAIM25/TAIM3/RandomImages/file03.jpg'  
      ... and 9 more  
    }  
    Folders: {  
      '/MATLAB Drive/TAIM25/TAIM3/RandomImages'  
    }  
    AlternateFileSystemRoots: {}  
    ReadSize: 1  
    Labels: {}  
    SupportedOutputFormats: ["png" "jpg" "jpeg" "tif" "tiff"]  
    DefaultOutputFormat: "png"  
    ReadFcn: @readDatastoreImage
```

# Create a Datastore for images

```
fname = imds.Files
```

```
img = readimage(imds,7);  
imshow(img)
```

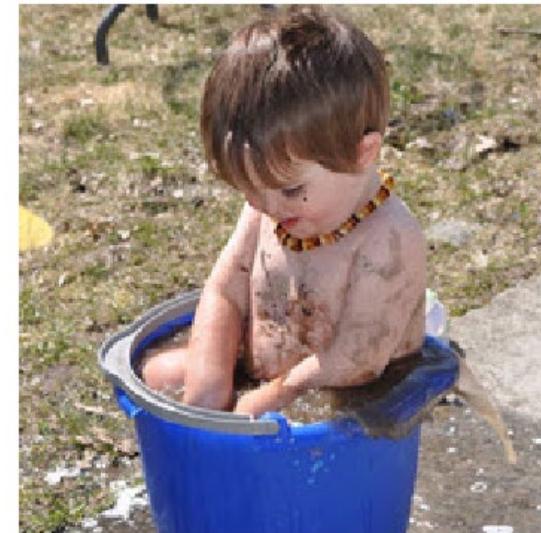
```
scores = minibatchpredict(net,imds)
```

```
preds = scores2label(scores,classes)
```

```
max(scores,[],2)
```

```
ans = 12x1 single column vector  
0.4867  
0.1700  
0.4020  
0.6952  
0.1643  
0.2414  
0.9361  
0.9991  
0.4436  
0.1995  
⋮
```

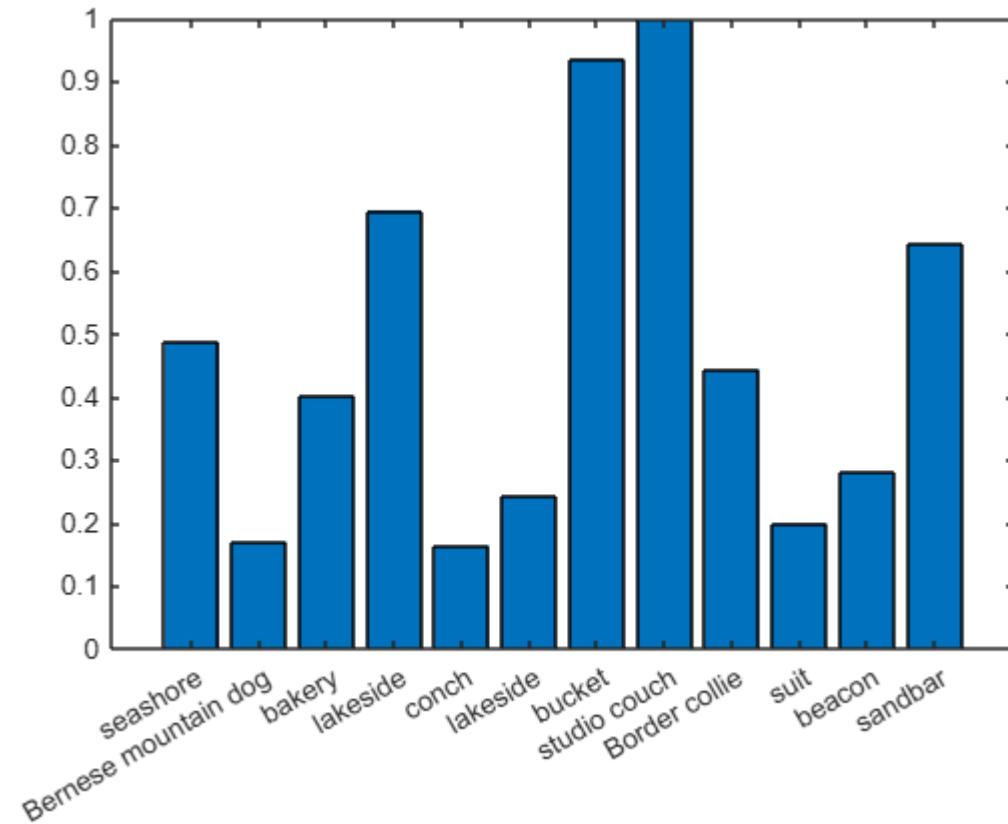
```
preds = 12x1 categorical  
seashore  
Bernese mountain d...  
bakery  
lakeside  
conch  
lakeside  
bucket  
studio couch  
Border collie  
suit  
⋮
```



# Create a Datastore for images

bar(max(scores,[],2))

xticklabels(preds)

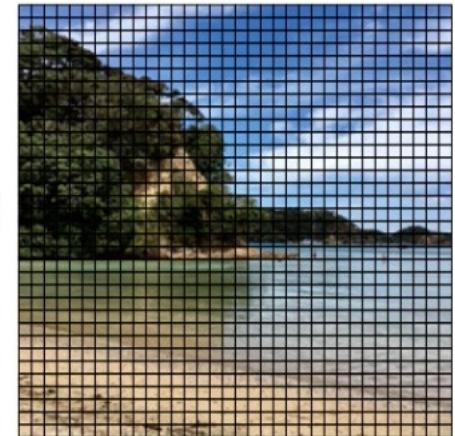


# Adjusting Input Images

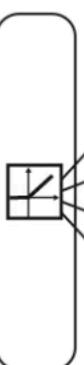
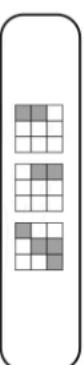
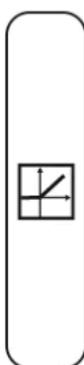
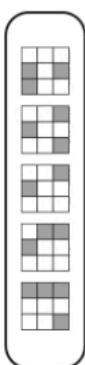
- What if the images does not have the “standard” characteristics expected by the pretrained NN?



227



227



Input

Convolution, Pooling, and ReLU

Fully Connected

Softmax

Color  
227 x 227

29

# Adjusting Input Images: going to code

## Adjust Input Images

```
% Let's see initial size of image file01.jpg
```

```
img = imread("file01.jpg");
```

```
imshow(img)
```

```
sz = size (img)
```

```
% On the other hand, let's see the main  
characteristics of the Network used:
```

```
summary(net)
```

```
% As initial size does not match, let's resize the image to the demands of
```

```
% the network
```

```
img = imresize (img, [224 224])
```

```
imshow (img)
```

```
% Finally, does it changes the result. Let's hope not!
```

```
s = minibatchpredict(net,img)
```

```
c = scores2label(s,classes)
```

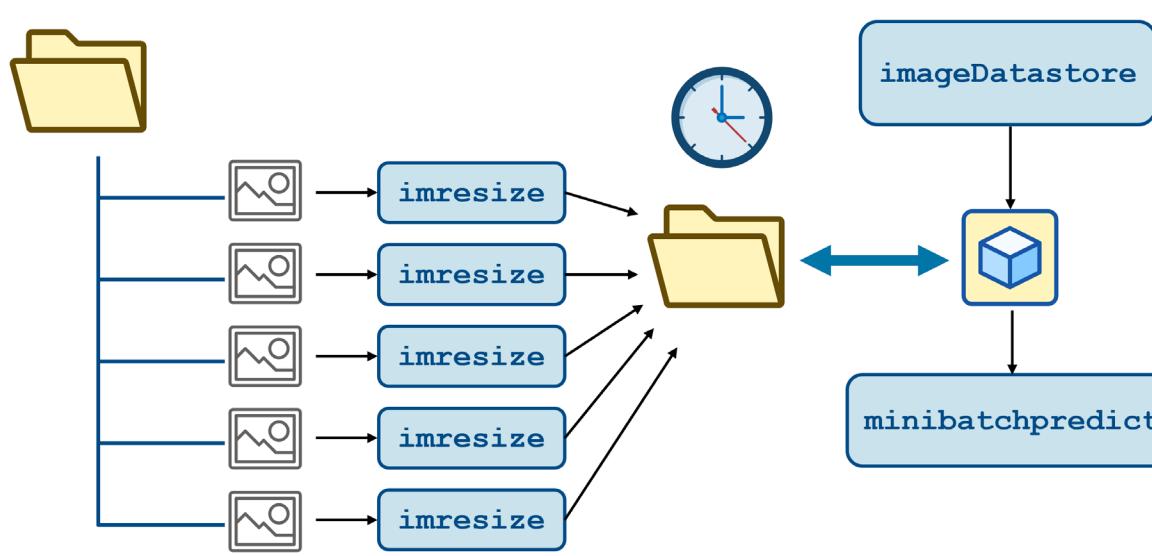
SZ = 1x3				
	1	2	3	
1	227	227	3	

```
Initialized: true  
Number of learnables: 6.9M  
Inputs:  
1 'data' 224x224x3 images
```

```
Initialized: true  
Number of learnables: 6.9M  
Inputs:  
1 'data' 224x224x3 images
```

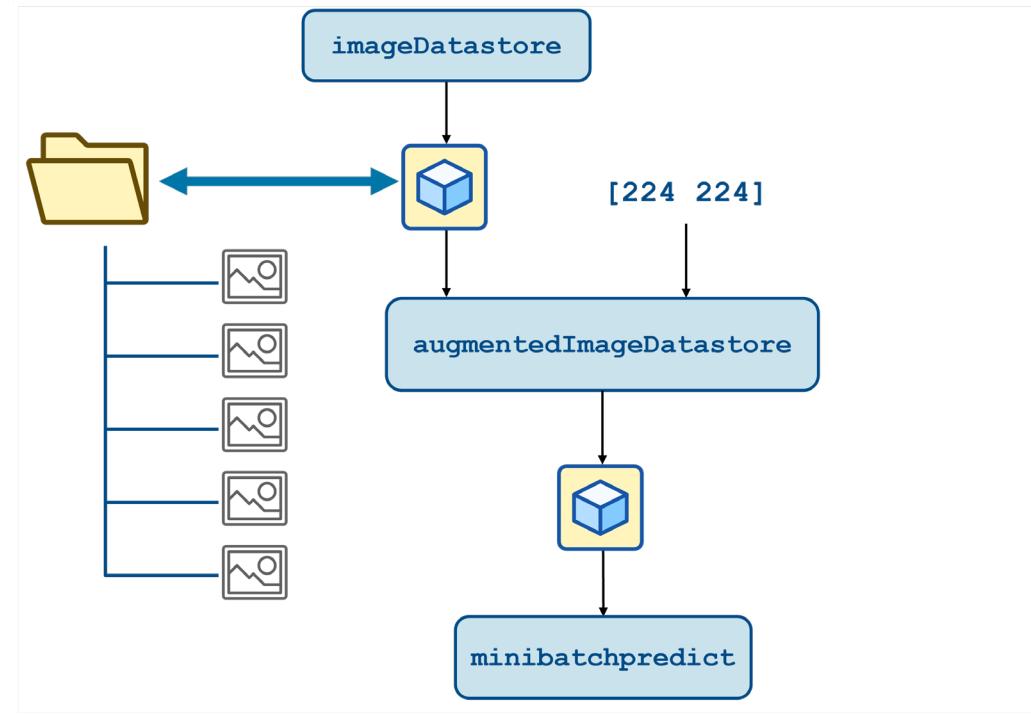
# Adjusting Input Images: Now over the “datastore”

- To use an image datastore, you also need to save the processed images to a folder.
- For large data sets, saving a duplicate of your files can take up a lot of space.



# Adjusting Input Images: Now over the “datastore”

- Alternatively, perform basic preprocessing with the “`augmentedImageDatastore`” function (it takes an image datastore and an image size as inputs and the datastore can be used as input to the `minibatchpredict` function).



# Adjusting Input Images: going to code

## Resize Images in a Datastore

% As indicated, we can create a completed "resized" image data store by:  
auds = augmentedImageDatastore([224 224],imds)

```
auds =
augmentedImageDatastore with properties:

    NumObservations: 12
        Files: {12x1 cell}
    AlternateFileSystemRoots: {}
        MiniBatchSize: 128
    DataAugmentation: 'none'
    ColorPreprocessing: 'none'
        OutputSize: [224 224]
    OutputSizeMode: 'resize'
    DispatchInBackground: 0
```

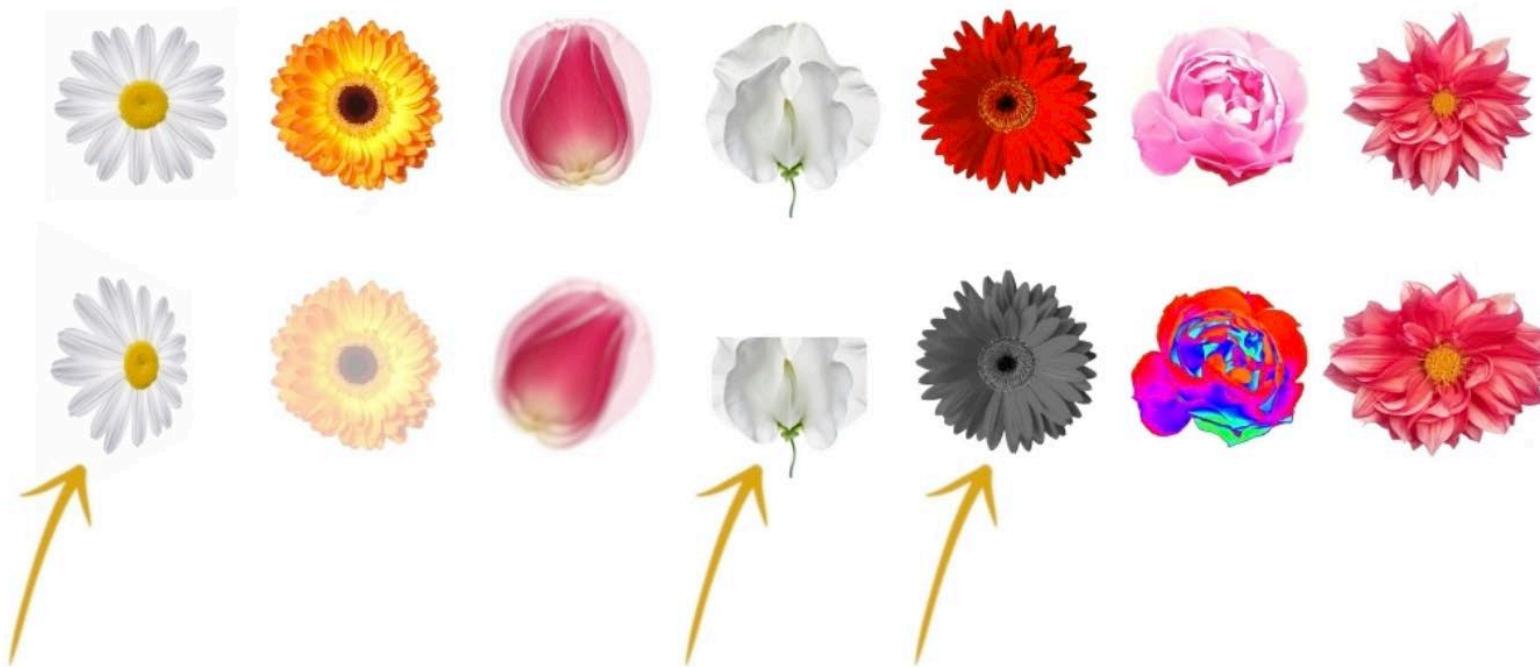
```
preds = 12x1 categorical
seashore
Bernese mountain d...
bakery
lakeside
chiton
lakeside
bucket
studio couch
Border collie
vestment
:
```

% So that we can process the whole resized collection of images directly:

```
scores = minibatchpredict(net,auds)
preds = scores2label(scores,classes)
```

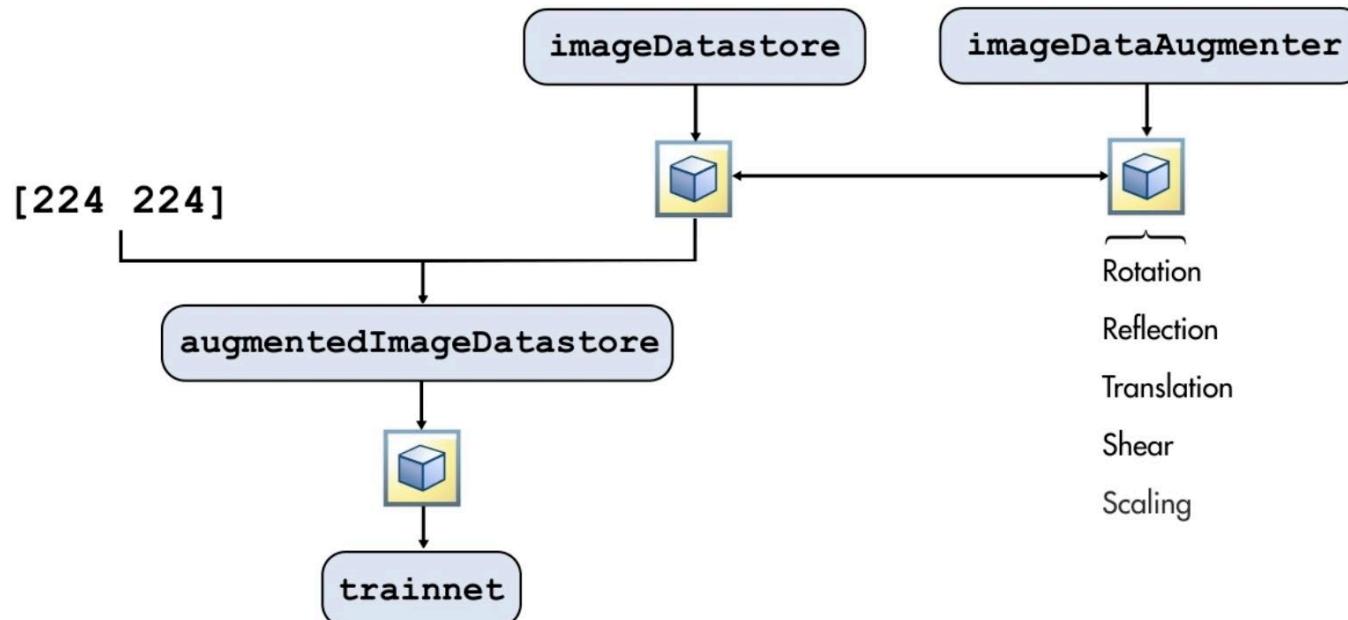
## Convert to RGB Images

- Maybe your training images may not be sufficiently varied.
- For example, to build a mobile phone app, you really want the classifier to be able to cope with imperfect images, odd angles, substandard lighting, not well-framed, or cropped, etc.



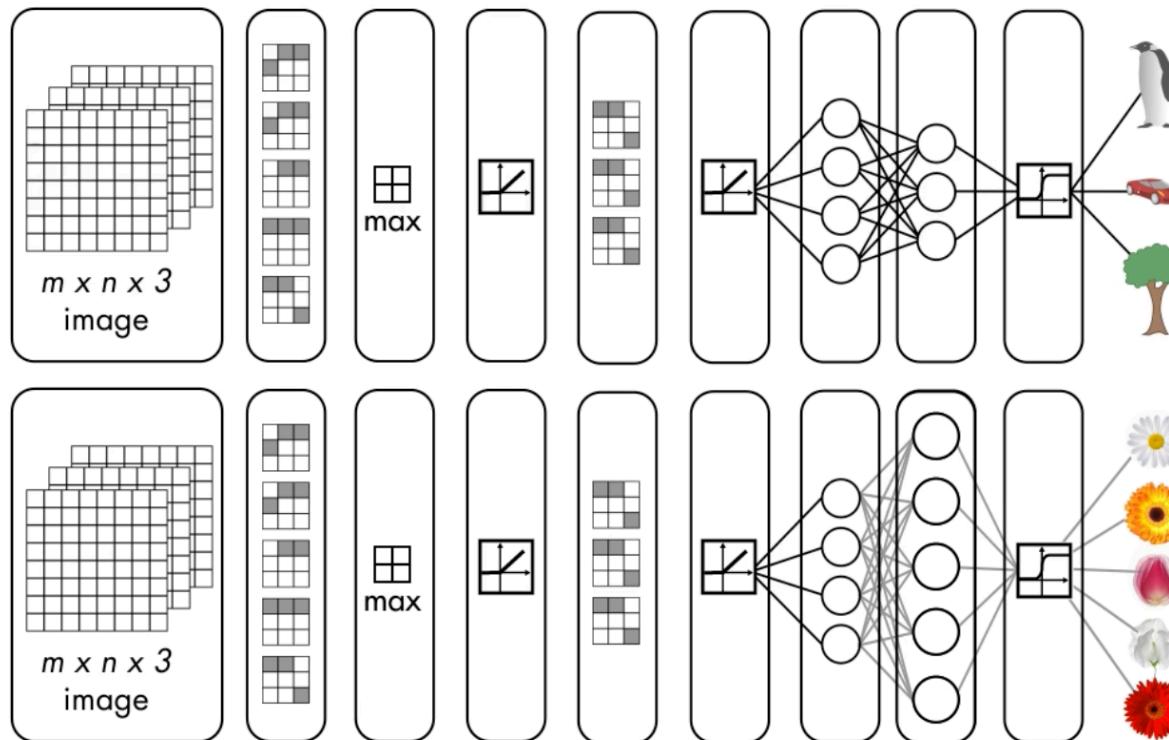
# Convert to RGB Images

- Use the “imageDataAugmenter” function to set up a set of transformations to your images to make an augmented training set.
- These transformations include: rotations, reflections, translations, shears, and scaling.
- You then pass the imageDataAugmenter variable to the “augmentedImageDatastore” function as an optional input.

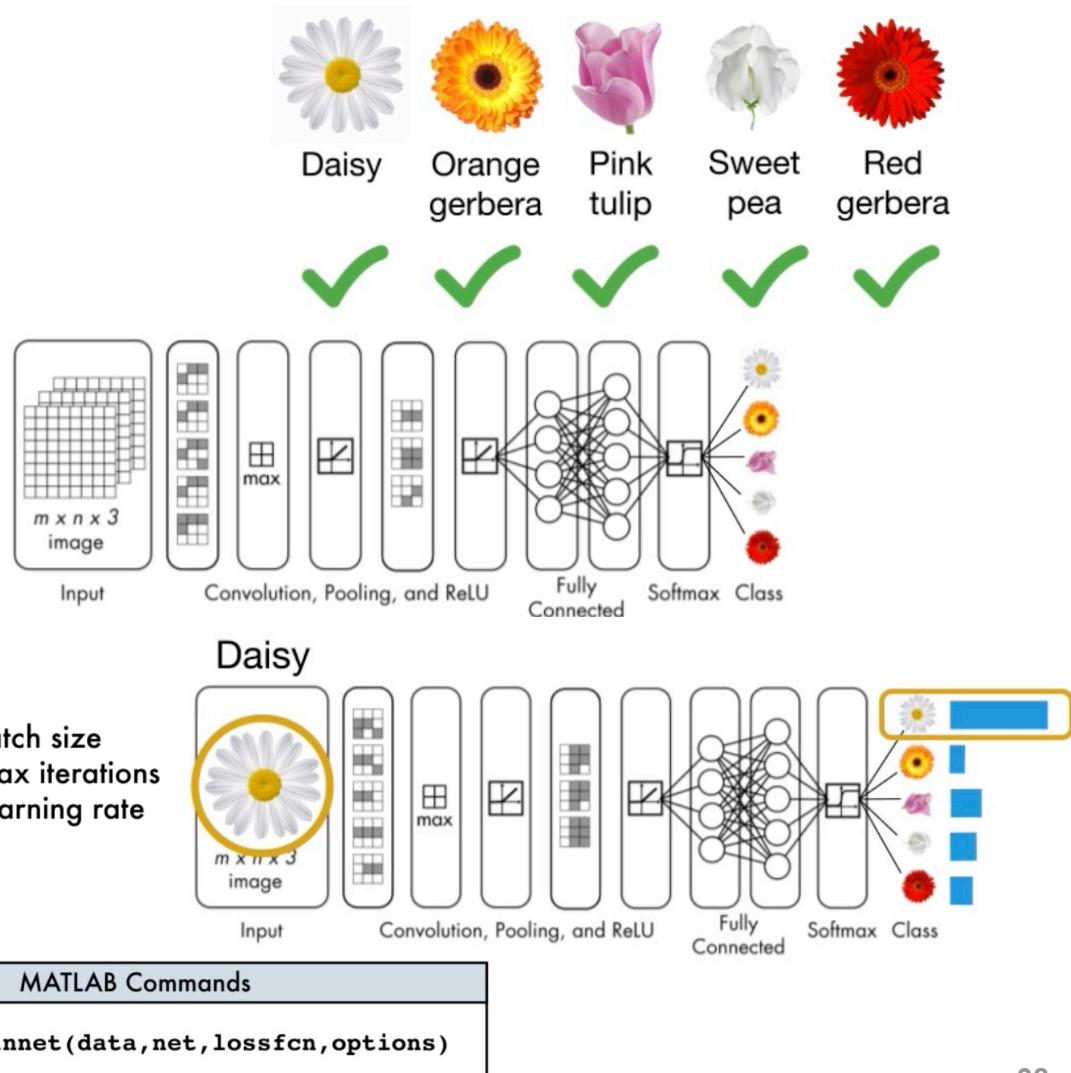
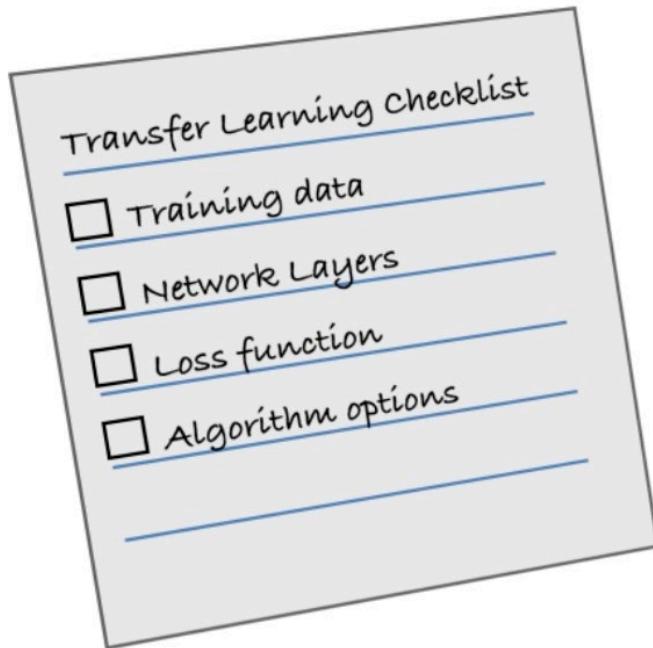


# What is Transfer Learning?

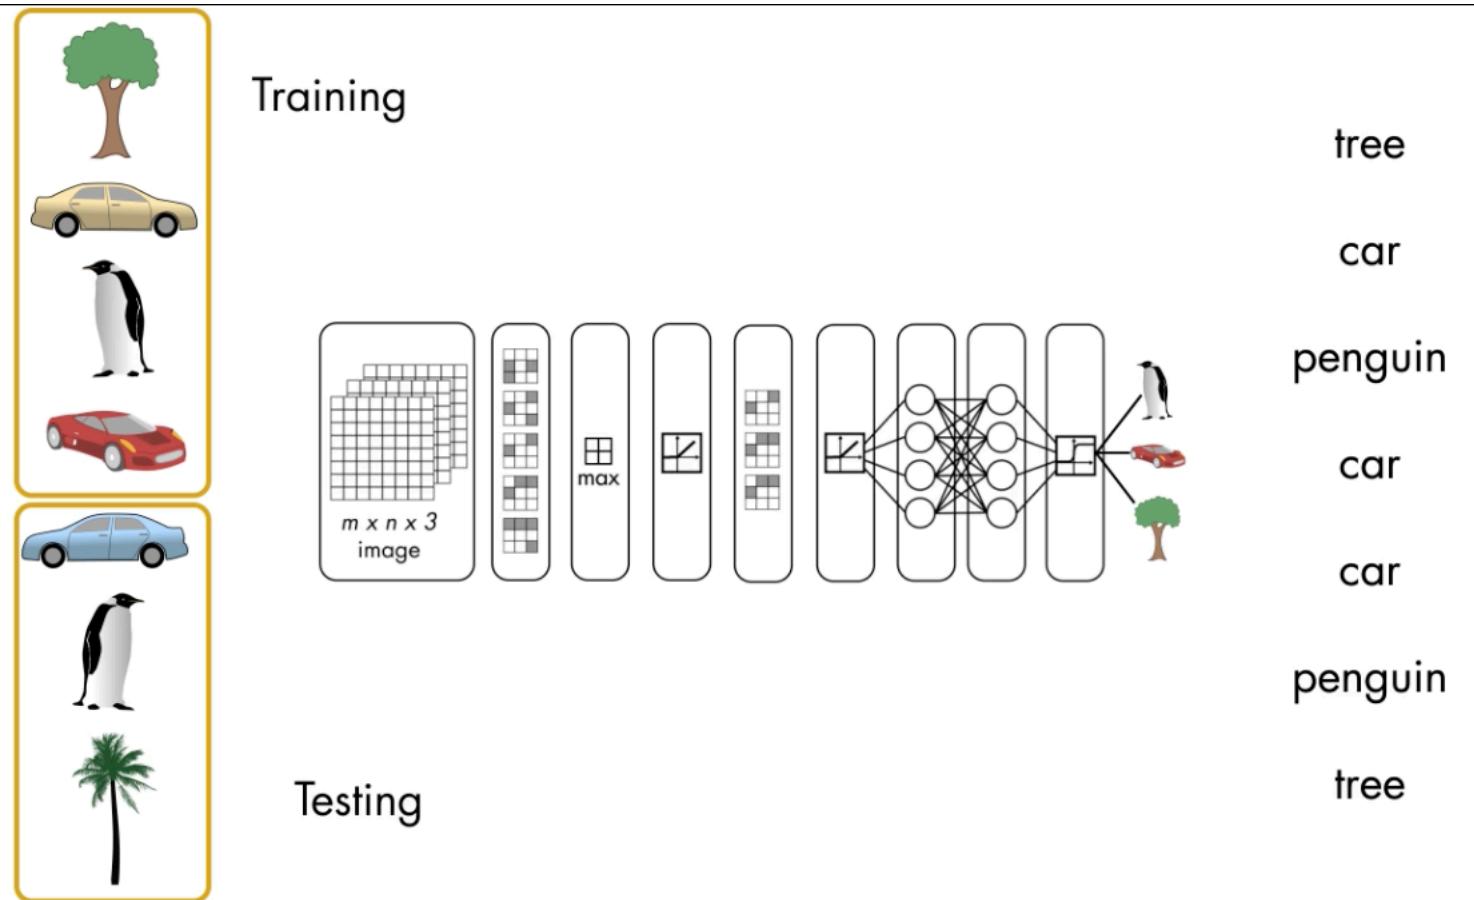
- Pretrained networks can classify a wide range of categories but may not cover your application's classes.
- Transfer learning involves taking a pretrained network, modifying it, and retraining it on new, specific data to fit the desired classification task.



# What do we need for Transfer Learning?

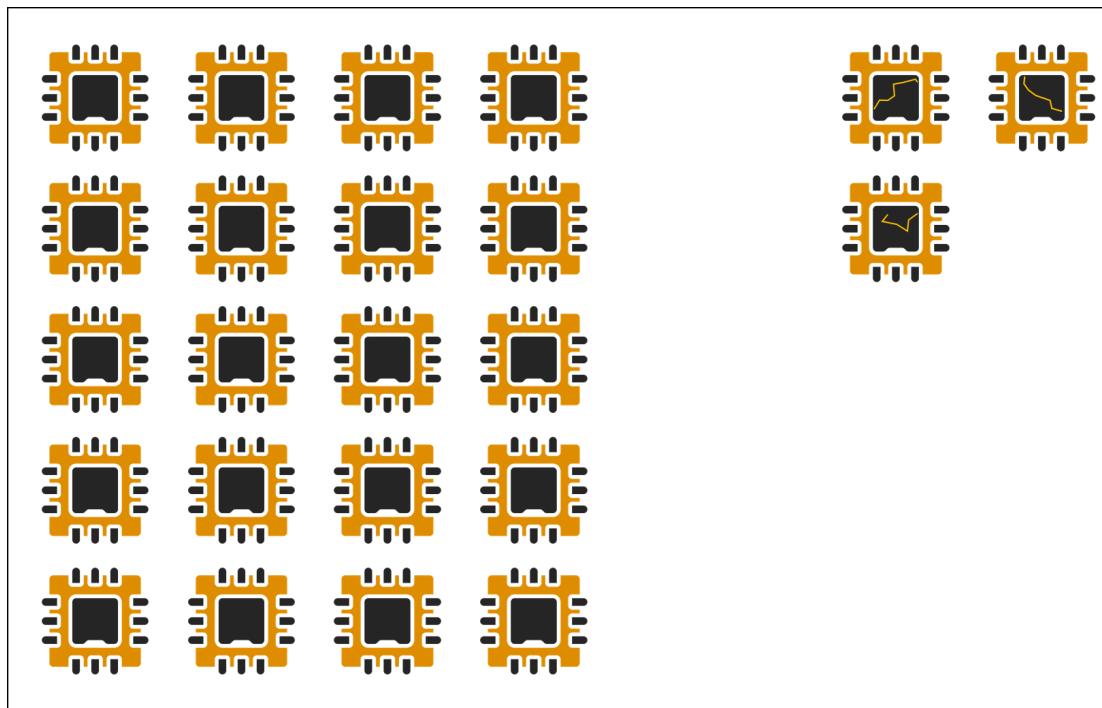


# As with Machine Learning, we need to slit data for Training and Testing



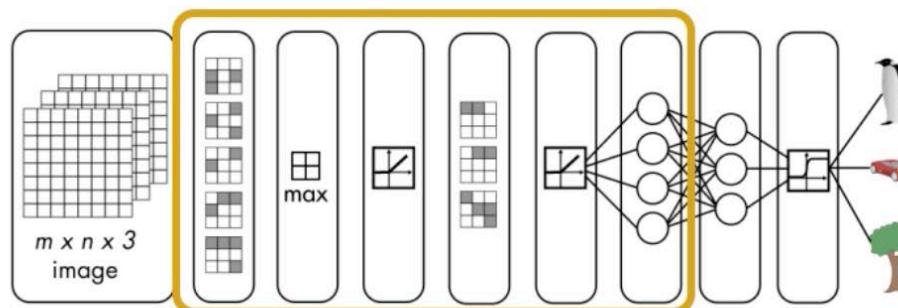
# Unbalanced Training Data

- Sometimes, you have many more images of one class than another.
- For example, if you try to detect defects, obtaining many images with no defect is usually easy, but obtaining images with defects is harder.
- In this case, this imbalance can bias the training.
- Splitting the data so that the training images have equal numbers from each class helps to avoid this bias.



# How to do transfer learning at a CNN

- Most of the layers of the pretrained network are: convolution, pooling, and rectified linear unit layers. These take the original input image and extract various features that can then be used for classification.



1	'data'	Image Input	224x224x3 images with 'zerocenter' normalization
2	'conv1-7x7_s2'	Convolution	64 7x7x3 convolutions with stride [2 2] and padding [3 3 3 3]
3	'conv1-relu_7x7'	ReLU	ReLU
4	'pool1-3x3_s2'	Max Pooling	3x3 max pooling with stride [2 2] and padding [0 1 0 1]
5	'pool1-norm1'	Cross Channel Normalization	cross channel normalization with 5 channels per element
6	'conv2-3x3_reduce'	Convolution	64 1x1x64 convolutions with stride [1 1] and padding [0 0 0 0]
7	'conv2-relu_3x3_reduce'	ReLU	ReLU
8	'conv2-3x3'	Convolution	192 3x3x64 convolutions with stride [1 1] and padding [1 1 1 1]
9	'conv2-relu_3x3'	ReLU	ReLU
10	'conv2-norm2'	Cross Channel Normalization	cross channel normalization with 5 channels per element
11	'pool2-3x3_s2'	Max Pooling	3x3 max pooling with stride [2 2] and padding [0 1 0 1]
12	'inception_3a-1x1'	Convolution	64 1x1x192 convolutions with stride [1 1] and padding [0 0 0 0]
13	'inception_3a-relu_1x1'	ReLU	ReLU
14	'inception_3a-3x3_reduce'	Convolution	96 1x1x192 convolutions with stride [1 1] and padding [0 0 0 0]
15	'inception_3a-relu_3x3_reduce'	ReLU	ReLU
16	'inception_3a-3x3'	Convolution	128 3x3x96 convolutions with stride [1 1] and padding [1 1 1 1]
17	'inception_3a-relu_3x3'	ReLU	ReLU
18	'inception_3a-5x5_reduce'	Convolution	16 1x1x192 convolutions with stride [1 1] and padding [0 0 0 0]
19	'inception_3a-relu_5x5_reduce'	ReLU	ReLU
20	'inception_3a-5x5'	Convolution	32 5x5x16 convolutions with stride [1 1] and padding [2 2 2 2]
21	'inception_3a-relu_5x5'	ReLU	ReLU
22	'inception_3a-pool'	Max Pooling	3x3 max pooling with stride [1 1] and padding [1 1 1 1]
23	'inception_3a-pool_proj'	Convolution	32 1x1x192 convolutions with stride [1 1] and padding [0 0 0 0]
24	'inception_3a-relu_pool_proj'	ReLU	ReLU

# How to do transfer learning at a CNN

- To perform transfer learning, you'll need to modify the network to suit your specific problem. But don't worry. Most of the network will stay the same. You typically change just the last few layers. For example, here's the end of

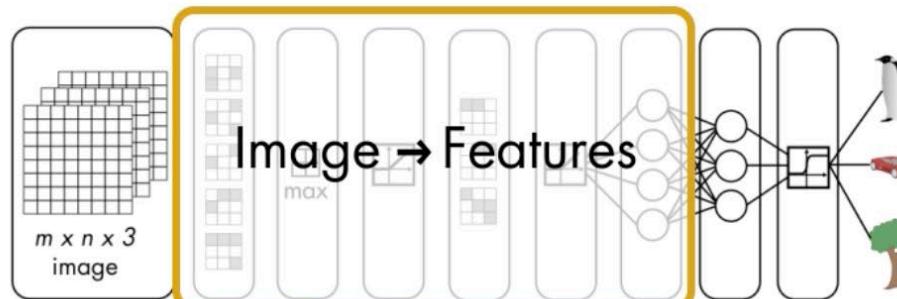
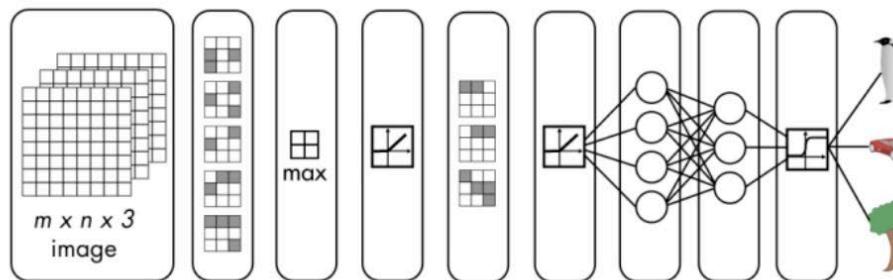


	Image Input	224x224x3 images with 'zerocenter' normalization
	Convolution	64 7x7x3 convolutions with stride [2 2] and padding [3 3 3 3]
1	'data'	ReLU
2	'conv1-7x7 s2'	Max Pooling
3	'conv1-relu_7x7'	Cross Channel Normalization
4	'pool1-3x3_s2'	Convolution
5	'pool1-norm1'	ReLU
6	'conv2-3x3_reduce'	Convolution
7	'conv2-relu_3x3_reduce'	ReLU
8	'conv2-3x3'	Convolution
9	'conv2-relu_3x3'	ReLU
10	'conv2-norm2'	Cross Channel Normalization
11	'pool2-3x3_s2'	Max Pooling
12	'inception_3a-1x1'	Convolution
13	'inception_3a-relu_1x1'	ReLU
14	'inception_3a-3x3_reduce'	Convolution
15	'inception_3a-relu_3x3_reduce'	ReLU
16	'inception_3a-3x3'	Convolution
17	'inception_3a-relu_3x3'	ReLU
18	'inception_3a-5x5_reduce'	Convolution
19	'inception_3a-relu_5x5_reduce'	ReLU
20	'inception_3a-5x5'	Convolution
21	'inception_3a-relu_5x5'	ReLU
22	'inception_3a-pool'	Max Pooling
23	'inception_3a-pool_proj'	Convolution
4	'inception_3a-relu pool proj'	ReLU

# How to do transfer learning at a CNN

- For transfer learning, you'll need to modify the network to your problem.
- Most of the network will stay the same. (You typically change just the last few layers.)

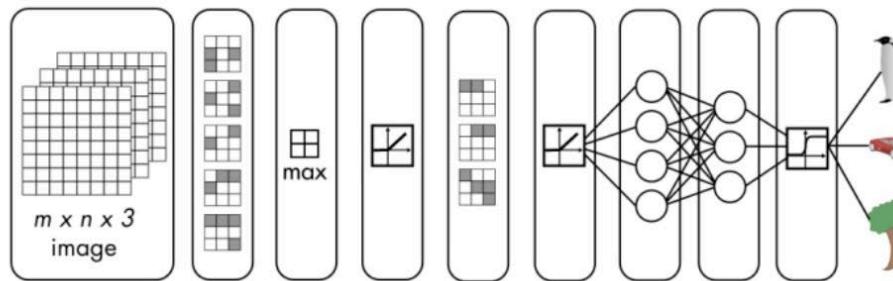


```
125  'inception_5a-output'          Depth concatenation
126  'inception_5b-1x1'            Convolution
127  'inception_5b-relu_1x1'       ReLU
128  'inception_5b-3x3_reduce'   Convolution
129  'inception_5b-relu_3x3_reduce' ReLU
130  'inception_5b-3x3'          Convolution
131  'inception_5b-relu_3x3'       ReLU
132  'inception_5b-5x5_reduce'   Convolution
133  'inception_5b-relu_5x5_reduce' ReLU
134  'inception_5b-5x5'          Convolution
135  'inception_5b-relu_5x5'       ReLU
136  'inception_5b-pool'          Max Pooling
137  'inception_5b-pool_proj'     Convolution
138  'inception_5b-relu_pool_proj' ReLU
139  'inception_5b-output'        Depth concatenation
140  'pool5-7x7_s1'              Global Average Pooling
141  'pool5-drop 7x7 s1'         Dropout
142  'loss3-classifier'          Fully Connected
143  'prob'                      Softmax
```

Depth concatenation of 4 inputs  
384 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
192 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
384 3x3x192 convolutions with stride [1 1] and padding [1 1 1 1]  
ReLU  
48 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
128 5x5x48 convolutions with stride [1 1] and padding [2 2 2 2]  
ReLU  
3x3 max pooling with stride [1 1] and padding [1 1 1 1]  
128 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
Depth concatenation of 4 inputs  
Global average pooling  
40% dropout  
1000 fully connected layer  
softmax

# How to do transfer learning at a CNN

- For example, the end of GoogLeNet, the 142nd layer is a fully connected layer with 1,000 neurons. (It takes the extracted features from the previous layers and maps them to the 1,000 output classes).

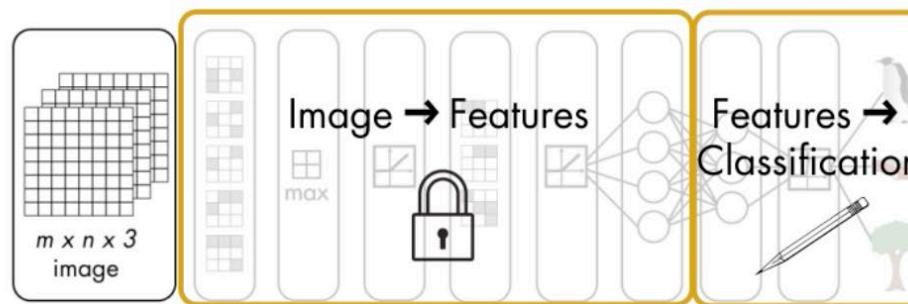


```
125  'inception_5a-output'          Depth concatenation
126  'inception_5b-1x1'            Convolution
127  'inception_5b-relu_1x1'       ReLU
128  'inception_5b-3x3_reduce'    Convolution
129  'inception_5b-relu_3x3_reduce' ReLU
130  'inception_5b-3x3'           Convolution
131  'inception_5b-relu_3x3'       ReLU
132  'inception_5b-5x5_reduce'    Convolution
133  'inception_5b-relu_5x5_reduce' ReLU
134  'inception_5b-5x5'           Convolution
135  'inception_5b-relu_5x5'       ReLU
136  'inception_5b-pool'          Max Pooling
137  'inception_5b-pool_proj'     Convolution
138  'inception_5b-relu_pool_proj' ReLU
139  'inception_5b-output'        Depth concatenation
140  'pool5-7x7_s1'              Global Average Pooling
141  'pool5-drop 7x7 s1'         Dropout
142  'loss3-classifier'          Fully Connected
143  'prob'                      Softmax
```

Depth concatenation of 4 inputs  
384 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
192 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
384 3x3x192 convolutions with stride [1 1] and padding [1 1 1 1]  
ReLU  
48 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
ReLU  
128 5x5x48 convolutions with stride [1 1] and padding [2 2 2 2]  
ReLU  
3x3 max pooling with stride [1 1] and padding [1 1 1 1]  
128 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
ReLU  
Depth concatenation of 4 inputs  
Global average pooling  
40% dropout  
1000 fully connected layer  
softmax

# How to do transfer learning at a CNN

- When changing these layers, the layer graph has the same feature extraction behavior as the pre-trained network.
- But it needs to be trained to map these features to your new image classes.



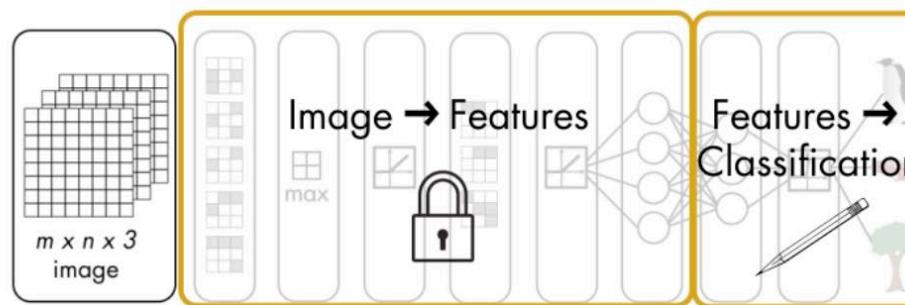
```
125 'inception_3a-output'          Depth concatenation
126 'inception_5b-1x1'            Convolution
127 'inception_5b-relu_1x1'       ReLU
128 'inception_5b-3x3_reduce'    Convolution
129 'inception_5b-relu_3x3_reduce' ReLU
130 'inception_5b-3x3'           Convolution
131 'inception_5b-relu_3x3'      ReLU
132 'inception_5b-5x5_reduce'    Convolution
133 'inception_5b-relu_5x5_reduce' ReLU
134 'inception_5b-5x5'           Convolution
135 'inception_5b-relu_5x5'      ReLU
136 'inception_5b-pool'          Max Pooling
137 'inception_5b-pool_proj'     Convolution
138 'inception_5b-relu_pool_proj' ReLU
139 'inception_5b-output'        Depth concatenation
140 'pool5-7x7_s1'              Global Average Pooling
141 'pool5-drop_7x7_s1'         Dropout
142 'loss3-classifier'          Fully Connected
143 'prob'                      Softmax
```

A lock icon is positioned next to the code line 136, indicating that the Max Pooling layer is frozen. A pencil icon is positioned next to the code line 142, indicating that the Fully Connected layer is trainable.

Depth concatenation of 4 inputs  
384 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
192 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
384 3x3x192 convolutions with stride [1 1] and padding [1 1 1 1]  
ReLU  
48 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
128 5x5x48 convolutions with stride [1 1] and padding [2 2 2 2]  
ReLU  
3x3 max pooling with stride [1 1] and padding [1 1 1 1]  
128 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
Depth concatenation of 4 inputs  
Global average pooling  
40% dropout  
1000 fully connected layer  
softmax

# How to do transfer learning at a CNN

- When you train with new data, the network will learn that mapping as well as refine the feature extraction to be slightly more specific to your application.



```
125 'inception_3a-output'          Depth concatenation
126 'inception_5b-1x1'            Convolution
127 'inception_5b-relu_1x1'       ReLU
128 'inception_5b-3x3_reduce'    Convolution
129 'inception_5b-relu_3x3_reduce' ReLU
130 'inception_5b-3x3'           Convolution
131 'inception_5b-relu_3x3'      ReLU
132 'inception_5b-5x5_reduce'    Convolution
133 'inception_5b-relu_5x5_reduce' ReLU
134 'inception_5b-5x5'           Convolution
135 'inception_5b-relu_5x5'      ReLU
136 'inception_5b-pool'          Max Pooling
137 'inception_5b-pool_proj'     Convolution
138 'inception_5b-relu_pool_proj' ReLU
139 'inception_5b-output'        Depth concatenation
140 'pool5-7x7_s1'              Global Average Pooling
141 'pool5-drop_7x7_s1'         Dropout
142 'loss3-classifier'          Fully Connected
143 'prob'                      Softmax
```

Depth concatenation of 4 inputs  
384 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
192 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
384 3x3x192 convolutions with stride [1 1] and padding [1 1 1 1]  
ReLU  
48 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
128 5x5x48 convolutions with stride [1 1] and padding [2 2 2 2]  
ReLU  
3x3 max pooling with stride [1 1] and padding [1 1 1 1]  
128 1x1x832 convolutions with stride [1 1] and padding [0 0 0 0]  
ReLU  
Depth concatenation of 4 inputs  
Global average pooling  
40% dropout  
1000 fully connected layer softmax



## Results of transfer learning a CNN

- When done, the newly trained network with layers like the input network, but with updated weights is given.
- A text display showing the progress of the training as:

```
newnet = trainnet(data,net,"crossentropy",options)
```

Iteration	Epoch	TimeElapsed	LearnRate	TrainingLoss	TrainingAccuracy
1	1	00:00:09	0.01	3.1155	26.562
10	2	00:01:13	0.01	0.6069	78.125
20	4	00:02:20	0.01	0.14191	92.969
30	6	00:03:26	0.01	0.034291	99.219
40	8	00:04:31	0.01	0.0052891	100
50	10	00:05:37	0.01	0.011322	100
60	12	00:06:42	0.01	0.0033386	100
70	14	00:07:47	0.01	0.00092616	100
80	16	00:08:53	0.01	0.0087101	99.219
90	18	00:09:55	0.01	0.0055897	100
100	20	00:10:57	0.01	0.0036221	100
110	22	00:12:00	0.01	0.0016095	100
120	24	00:13:05	0.01	0.00037103	100
130	26	00:14:10	0.01	0.00018275	100
140	28	00:15:14	0.01	0.00015373	100
150	30	00:16:21	0.01	0.0001372	100

Training stopped: Max epochs completed

## Results of transfer learning a CNN

- Accuracy is the percentage of the training images that the network classifies correctly. (Expected increasing during your training).
- However, accuracy doesn't measure how confident the network is about each classification.

```
newnet = trainnet(data,net,"crossentropy",options)
```

Iteration	Epoch	TimeElapsed	LearnRate	TrainingLoss	TrainingAccuracy
1	1	00:00:09	0.01	3.1155	26.562
10	2	00:01:13	0.01	0.6069	78.125
20	4	00:02:20	0.01	0.14191	92.969
30	6	00:03:26	0.01	0.034291	99.219
40	8	00:04:31	0.01	0.0052891	100
50	10	00:05:37	0.01	0.011322	100
60	12	00:06:42	0.01	0.0033386	100
70	14	00:07:47	0.01	0.00092616	100
80	16	00:08:53	0.01	0.0087101	99.219
90	18	00:09:55	0.01	0.0055897	100
100	20	00:10:57	0.01	0.0036221	100
110	22	00:12:00	0.01	0.0016095	100
120	24	00:13:05	0.01	0.00037103	100
130	26	00:14:10	0.01	0.00018275	100
140	28	00:15:14	0.01	0.00015373	100
150	30	00:16:21	0.01	0.0001372	100

Training stopped: Max epochs completed

## Results of transfer learning a CNN

- The loss measures how far from a perfect prediction the network is (totaled over the set of training images).
- This should decrease towards zero as the training proceeds.

```
newnet = trainnet(data,net,"crossentropy",options)
```

Iteration	Epoch	TimeElapsed	LearnRate	TrainingLoss	TrainingAccuracy
1	1	00:00:09	0.01	3.1155	26.562
10	2	00:01:13	0.01	0.6069	78.125
20	4	00:02:20	0.01	0.14191	92.969
30	6	00:03:26	0.01	0.034291	99.219
40	8	00:04:31	0.01	0.0052891	100
50	10	00:05:37	0.01	0.011322	100
60	12	00:06:42	0.01	0.0033386	100
70	14	00:07:47	0.01	0.00092616	100
80	16	00:08:53	0.01	0.0087101	99.219
90	18	00:09:55	0.01	0.0055897	100
100	20	00:10:57	0.01	0.0036221	100
110	22	00:12:00	0.01	0.0016095	100
120	24	00:13:05	0.01	0.00037103	100
130	26	00:14:10	0.01	0.00018275	100
140	28	00:15:14	0.01	0.00015373	100
150	30	00:16:21	0.01	0.0001372	100

Training stopped: Max epochs completed

## Results of transfer learning a CNN

- You may often see the accuracy plateau at its maximum value, while the loss value continues to decrease.
- This is good. It means that the network is still getting better at differentiating between the categories, even if the final classification may not be changing.

```
newnet = trainnet(data,net,"crossentropy",options)
```

Iteration	Epoch	TimeElapsed	LearnRate	TrainingLoss	TrainingAccuracy
1	1	00:00:09	0.01	3.1155	26.562
10	2	00:01:13	0.01	0.6069	78.125
20	4	00:02:20	0.01	0.14191	92.969
30	6	00:03:26	0.01	0.034291	99.219
40	8	00:04:31	0.01	0.0052891	100
50	10	00:05:37	0.01	0.011322	100
60	12	00:06:42	0.01	0.0033386	100
70	14	00:07:47	0.01	0.00092616	100
80	16	00:08:53	0.01	0.0087101	99.219
90	18	00:09:55	0.01	0.0055897	100
100	20	00:10:57	0.01	0.0036221	100
110	22	00:12:00	0.01	0.0016095	100
120	24	00:13:05	0.01	0.00037103	100
130	26	00:14:10	0.01	0.00018275	100
140	28	00:15:14	0.01	0.00015373	100
150	30	00:16:21	0.01	0.0001372	100

Training stopped: Max epochs completed

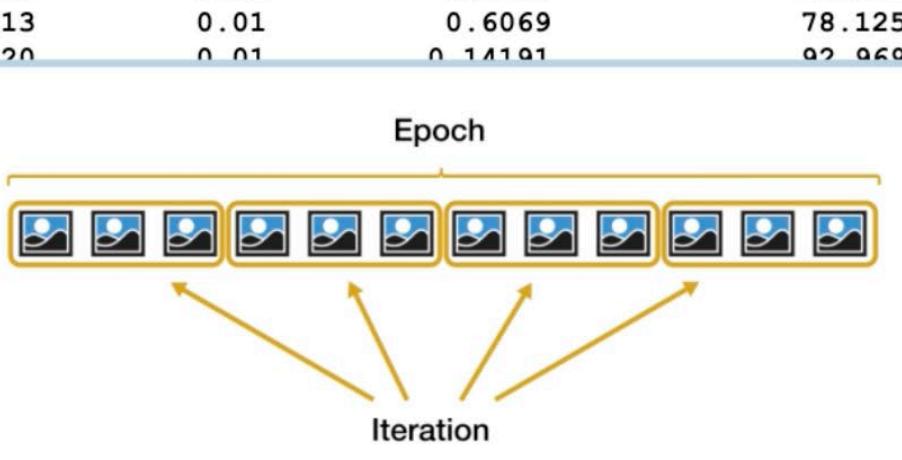
## Results of transfer learning a CNN

- At end of the table, it can be seen what caused training to stop.
- Usually due to the number of times the network has processed the training data.

```
newnet = trainnet(data,net,"crossentropy",options)
```

Iteration	Epoch	TimeElapsed	LearnRate	TrainingLoss	TrainingAccuracy
1	1	00:00:09	0.01	3.1155	26.562
10	2	00:01:13	0.01	0.6069	78.125
20	4	00:02:20	0.01	0.14191	92.969
30	6	00:03:29	0.01	0.00015373	100
40	8	00:04:37	0.01	0.0001372	100
50	10	00:05:45	0.01		
60	12	00:06:53	0.01		
70	14	00:07:59	0.01		
80	16	00:08:57	0.01		
90	18	00:09:55	0.01		
100	20	00:10:53	0.01		
110	22	00:11:51	0.01		
120	24	00:12:49	0.01		
130	26	00:13:47	0.01		
140	28	00:15:14	0.01	0.00015373	100
150	30	00:16:21	0.01	0.0001372	100

Training stopped: Max epochs completed



## Results of transfer learning a CNN

- At each iteration, the network training uses a subset of the training images, known as a mini-batch, to update the weights. Once the network has trained on the whole training data set, it has completed one epoch. To train the network for more or less time, adjust the maximum epochs option.

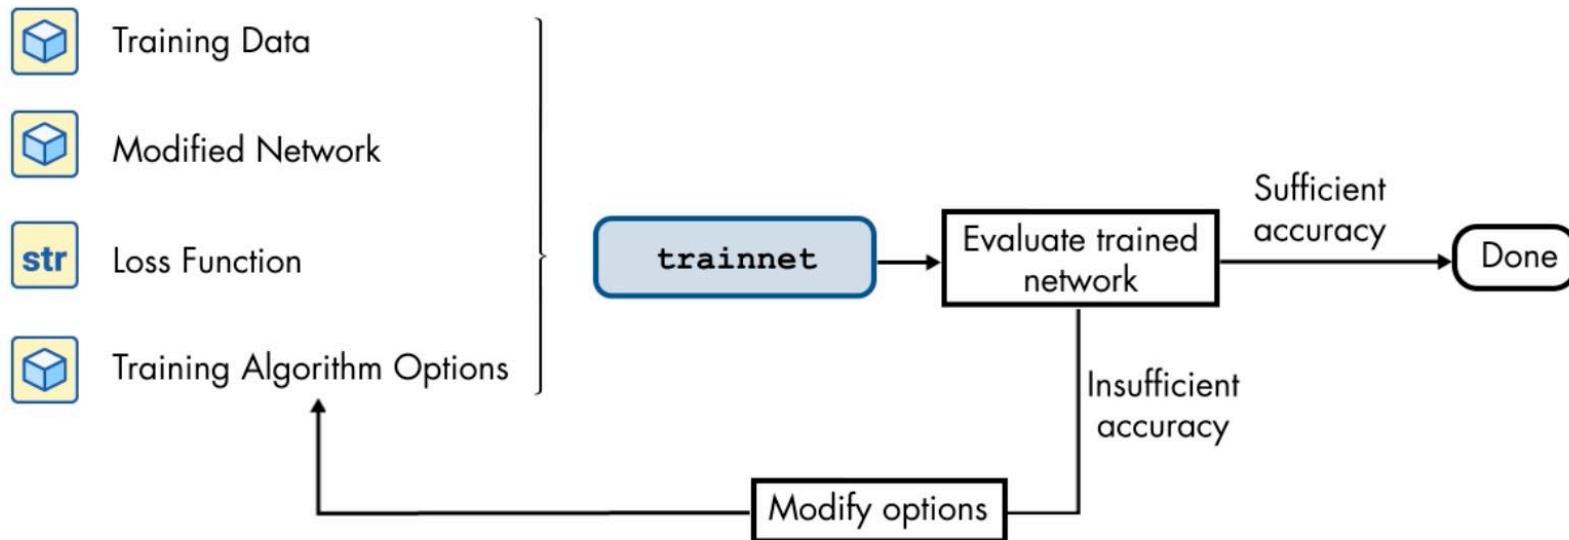
```
newnet = trainnet(data,net,"crossentropy",options)
```

Iteration	Epoch	TimeElapsed	LearnRate	TrainingLoss	TrainingAccuracy
1	1	00:00:09	0.01	3.1155	26.562
10	2	00:01:13	0.01	0.6069	78.125
20	4	00:02:20	0.01	0.14191	92.969
30	6	00:03:29	0.01	0.00015373	100
40	8	00:04:37	0.01	0.0001372	100
50	10	00:05:45	0.01		
60	12	00:06:53	0.01		
70	14	00:07:59	0.01		
80	16	00:08:57	0.01		
90	18	00:09:55	0.01		
100	20	00:10:53	0.01		
110	22	00:11:51	0.01		
120	24	00:12:49	0.01		
130	26	00:13:47	0.01		
140	28	00:15:14	0.01	0.00015373	100
150	30	00:16:21	0.01	0.0001372	100

Training stopped: Max epochs completed

# What if the accuracy is not enough?

- If the performance is good enough, then congratulations,
- Often the results you get when you first train the network aren't as good as you need.
- What do you do?

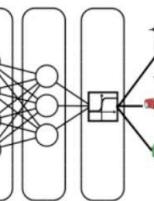
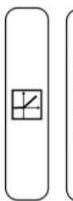
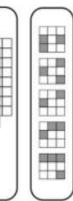
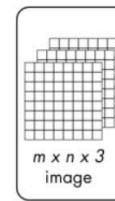
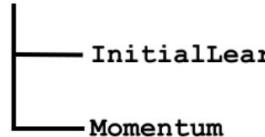


# What if the accuracy is not enough?

- If the performance is good enough, then congratulations,
- Often the results you get when you first train the network aren't as good as you need.
- What do you do?



Training Algorithm Options

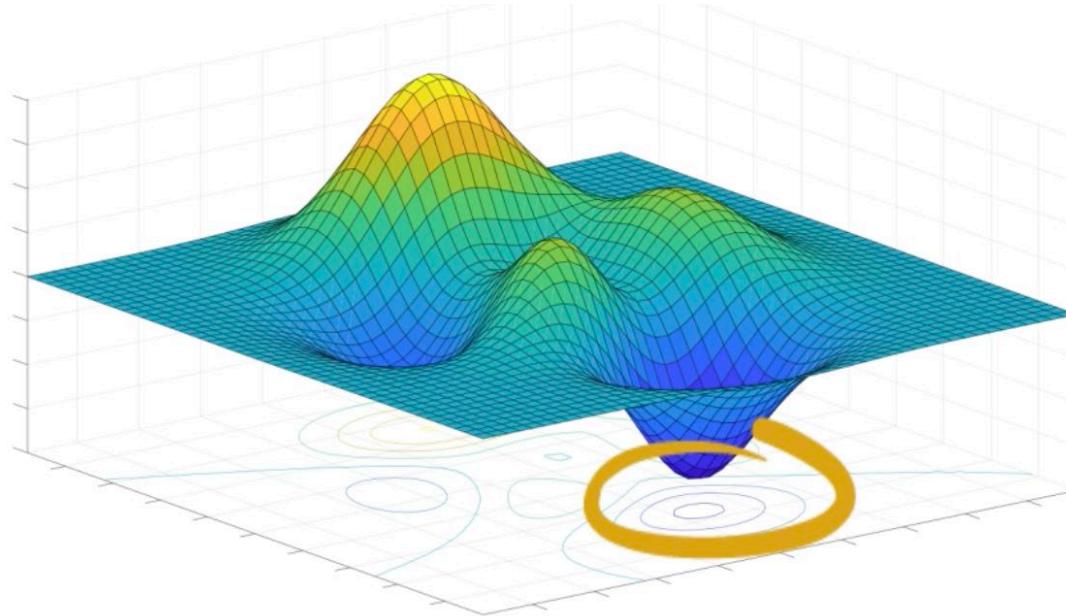


Prediction	Reality
65% tree	penguin
90% car	car
70% penguin	penguin
55% car	tree
80% car	car
50% penguin	car
95% tree	tree
75% tree	penguin

Loss = Difference

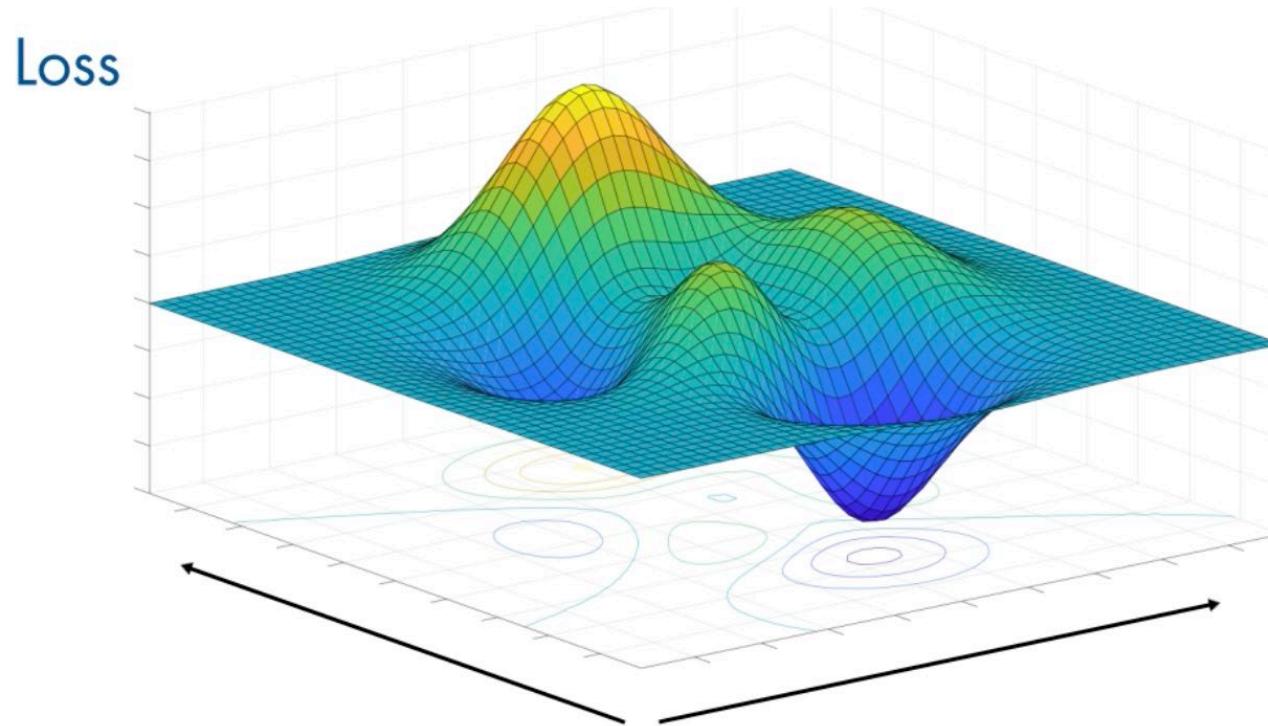
## How to visualize the training process?

- Common way to visualize the training process is to imagine that you're trying to find the bottom of a valley.



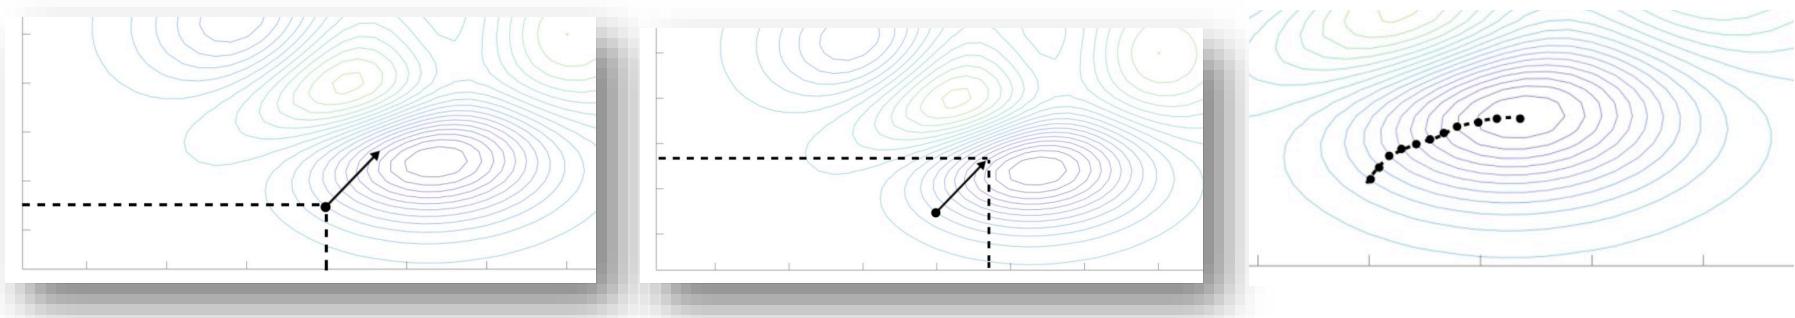
## How to visualize the training process?

- The height of the ground is the value of the loss function.
- Your north-south location is one weight value, and your east-west location is another. Of course, real networks have thousands of weights, not just two, but the principle is the same.



## How to visualize the training process?

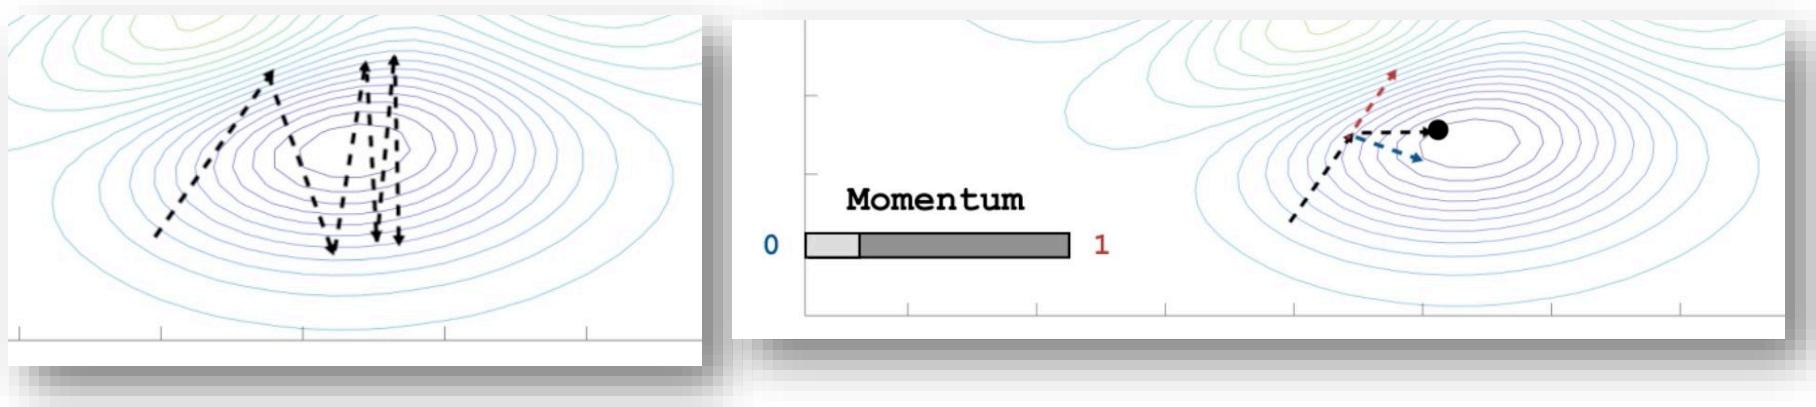
- The gradient descent algorithm works by looking at the height of the ground at a few places around your current position and using that to figure out the gradient or direction that the ground slopes down the steepest.



- Then you simply take a step in that direction, which brings you to a new set of weight values.
- But how big a step do you take?
- The size of that step is the learning rate.
- If the learning rate is too small, you'll need to take a lot of steps to get to the bottom, and that takes time.
- But if the learning rate is too big, you might end up jumping around all over the place.

## How to visualize the training process?

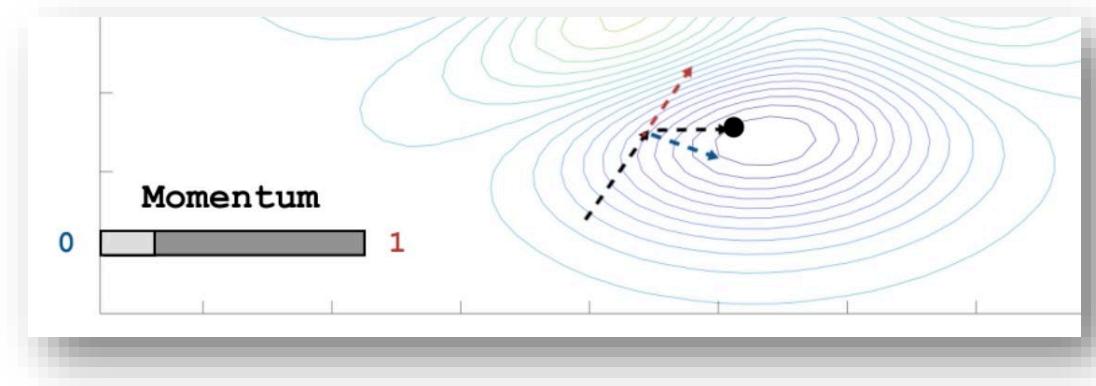
- If your weights have blown up and your network is giving terrible performance or even not working at all, the first thing you should do is divide your learning rate by 10 and try again.



- Gradient descent with momentum tries to stop the jumping around without having to take really small steps.
- You start off going straight down the direction of the gradient

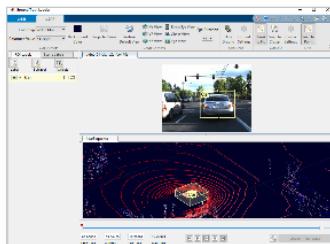
## How to visualize the training process?

- Instead, you turn toward that direction, similar to the way you would if you were a mass traveling in your original direction but now being pulled in the new direction by a force.
- Your next step is therefore in a direction that's a weighted average of the previous direction and the new direction. That weighting is set by the momentum option.



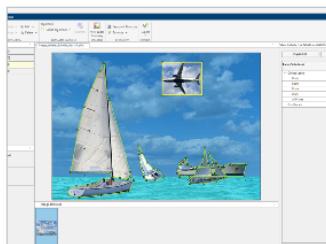
# More resources (if you want)

- You have also examples at:
- [https://uk.mathworks.com/help/deeplearning/examples.html?category=applications&exampleproduct=all&newonly=&s\\_tid=CRUX\\_lftnav](https://uk.mathworks.com/help/deeplearning/examples.html?category=applications&exampleproduct=all&newonly=&s_tid=CRUX_lftnav)
- And Apps for different Applications
- [https://uk.mathworks.com/help/deeplearning/referencelist.html?type=app&s\\_id=CRUX\\_topnav&category=applications](https://uk.mathworks.com/help/deeplearning/referencelist.html?type=app&s_id=CRUX_topnav&category=applications)



#### Ground Truth Labeler

The Ground Truth Labeler app enables you to label ground truth data in multiple videos, image sequences, or lidar point clouds.



#### Image Labeler

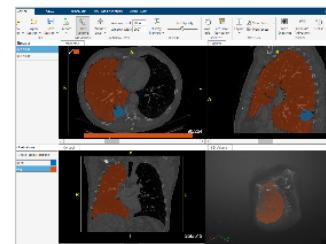
The Image Labeler app enables you to label ground truth data in a collection of images. Using the app, you can:



#### Lidar Labeler

The Lidar Labeler app enables you to label objects in a point cloud or a point cloud sequence. The app reads point cloud data from PLY, PCAP, LA...

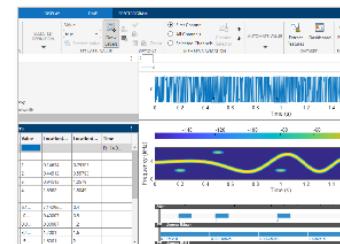
Since R2020b



#### Medical Image Labeler

The Medical Image Labeler app enables you to label ground truth data in medical images. Using the app, you can:

Since R2022b



#### Signal Labeler

The Signal Labeler app is an interactive tool that enables you to label signals for analysis or for use in machine learning and deep learning...